




AN ACCELERATED SYSTEM IS DIFFERENT

Steve Oberlin @ GTCDC 2018

WHAT IS THE PROBLEM?

THE CPU IS OUT OF GAS


No Really Running on Fumes



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten New plot and data collected for 2010-2015 by K. Rupp

PRICE/PERFORMANCE ALREADY CONSTRAINED

CPU trends



- CPU evolution is not able to cope with the increasing demand of performance
- Depending on the application, GPUs can provide better performance and energy efficiency

SENSE OF URGENCY?

Dennard scaling ended 13 years ago

Single threaded performance is capped at ~ 2.5 GHz

Moore's law is slowing



Transistor density increasing at a slower rate

The economics have changed

Fewer transistors per generation cost more to produce, and don't run any faster

What's Your Plan?

AN ACCELERATED SYSTEM CAN MAINTAIN OR EXCEED PREVIOUS PERFORMANCE TRAJECTORY





Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten New plot and data collected for 2010-2015 by K. Rupp

BUT NOT ALL APPS can move from the **BLUE line** to the **GREEN line**
And their “mileage” will vary

THE SOLUTION

NEW ALGORITHMS AND METHODS ARE EMERGING TO DEFINE A NEW NORMAL IN HPC



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten New plot and data collected for 2010-2015 by K. Rupp

THE KEY DIFFERENCES

Unaccelerated systems

Have homogenous nodes, where all nodes are the same

Max throughput with minimum variance regardless of workload

All the apps are modeling/simulation

Most apps are batch

Most modeling/simulation apps are double precision

Accelerated systems

Have heterogenous nodes, where nodes are tuned to the workload


Max throughput and maximize performance variance based on workload

Wide range of apps: modeling/simulation to data science

A mix of batch and interactive jobs

Precision selected to optimize performance

THE PATH FORWARD



+




CPU + Accelerator

$$\begin{aligned}\nabla \cdot \mathbf{E} &= \rho / \epsilon_0 & \nabla \cdot \mathbf{B} &= 0 \\ \nabla \times \mathbf{E} &= -\dot{\mathbf{B}} & \nabla \times \mathbf{B} &= \mu_0 \mathbf{j} + \mu_0 \epsilon_0 \dot{\mathbf{E}}\end{aligned}$$



Simulation + AI

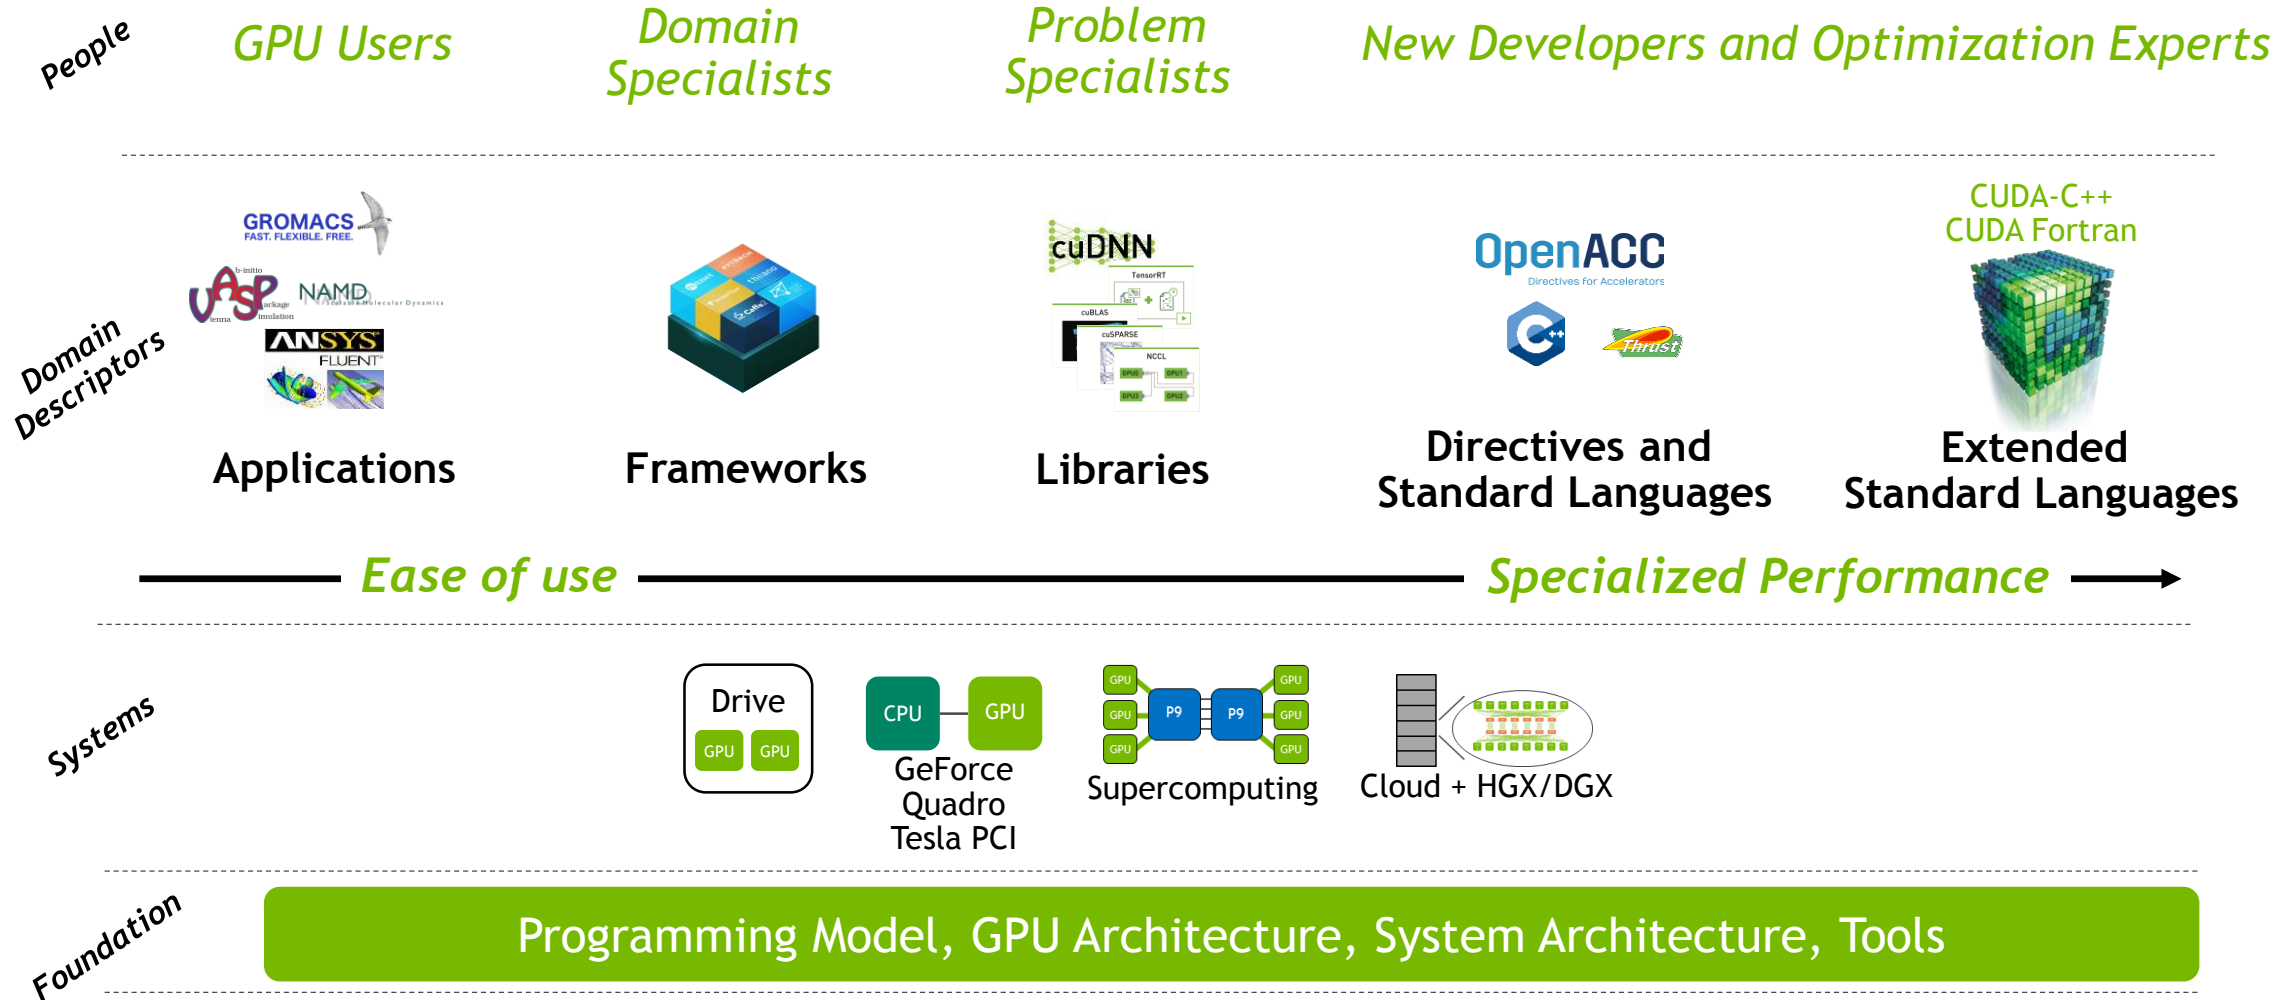


FP64 + Multi-Precision



Full-stack Optimization

CUDA PLATFORM

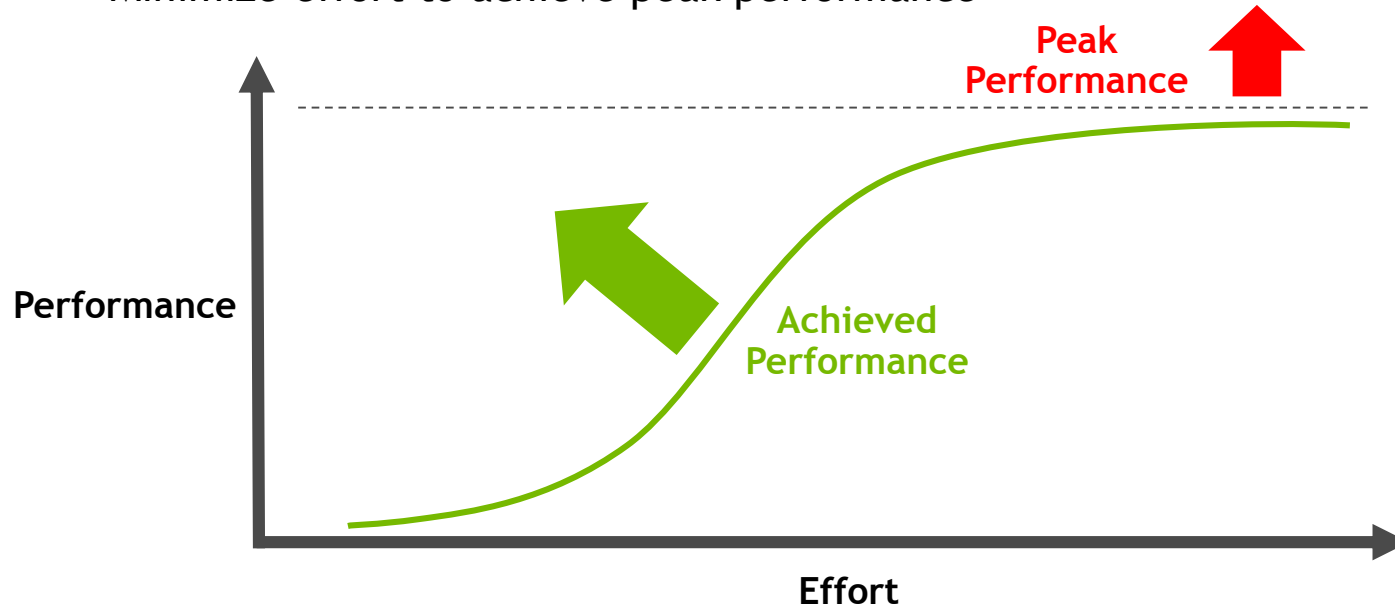


CUDA PLATFORM GOALS

Every **layer** of the CUDA Platform has two fundamental goals:

Maximize achievable peak performance

Minimize effort to achieve peak performance



CUDA Platform Layers

Applications

Frameworks & DSLs

Libraries

Directives and
Standard Languages

Extended
Standard Languages


System Architecture

GPU Architecture

AI: A NEW TOOL FOR SCIENCE



The SKA1 Square Kilometre Array radio telescope will generate more than an Exabyte of data every day.




The CERN large Hadron collider's High Luminosity upgrade will result in a 10X increase in data volume.



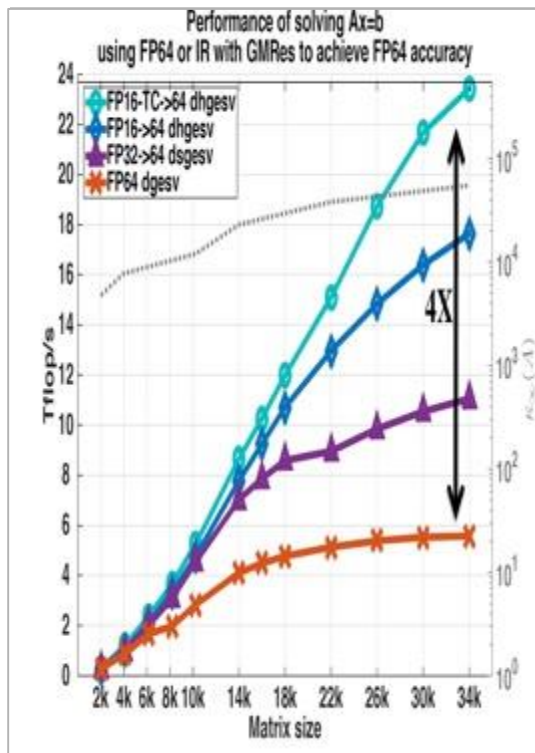
The 500 MW ITER fusion experiment will provide a 30X increase in output power over the largest previous experiment.

TENSOR CORES FOR SCIENCE


Multi-precision computing




FP64+ MULTI-PRECISION



MAGMA Library with FP16




PLASMA FUSION
APPLICATION



EARTHQUAKE SIMULATION

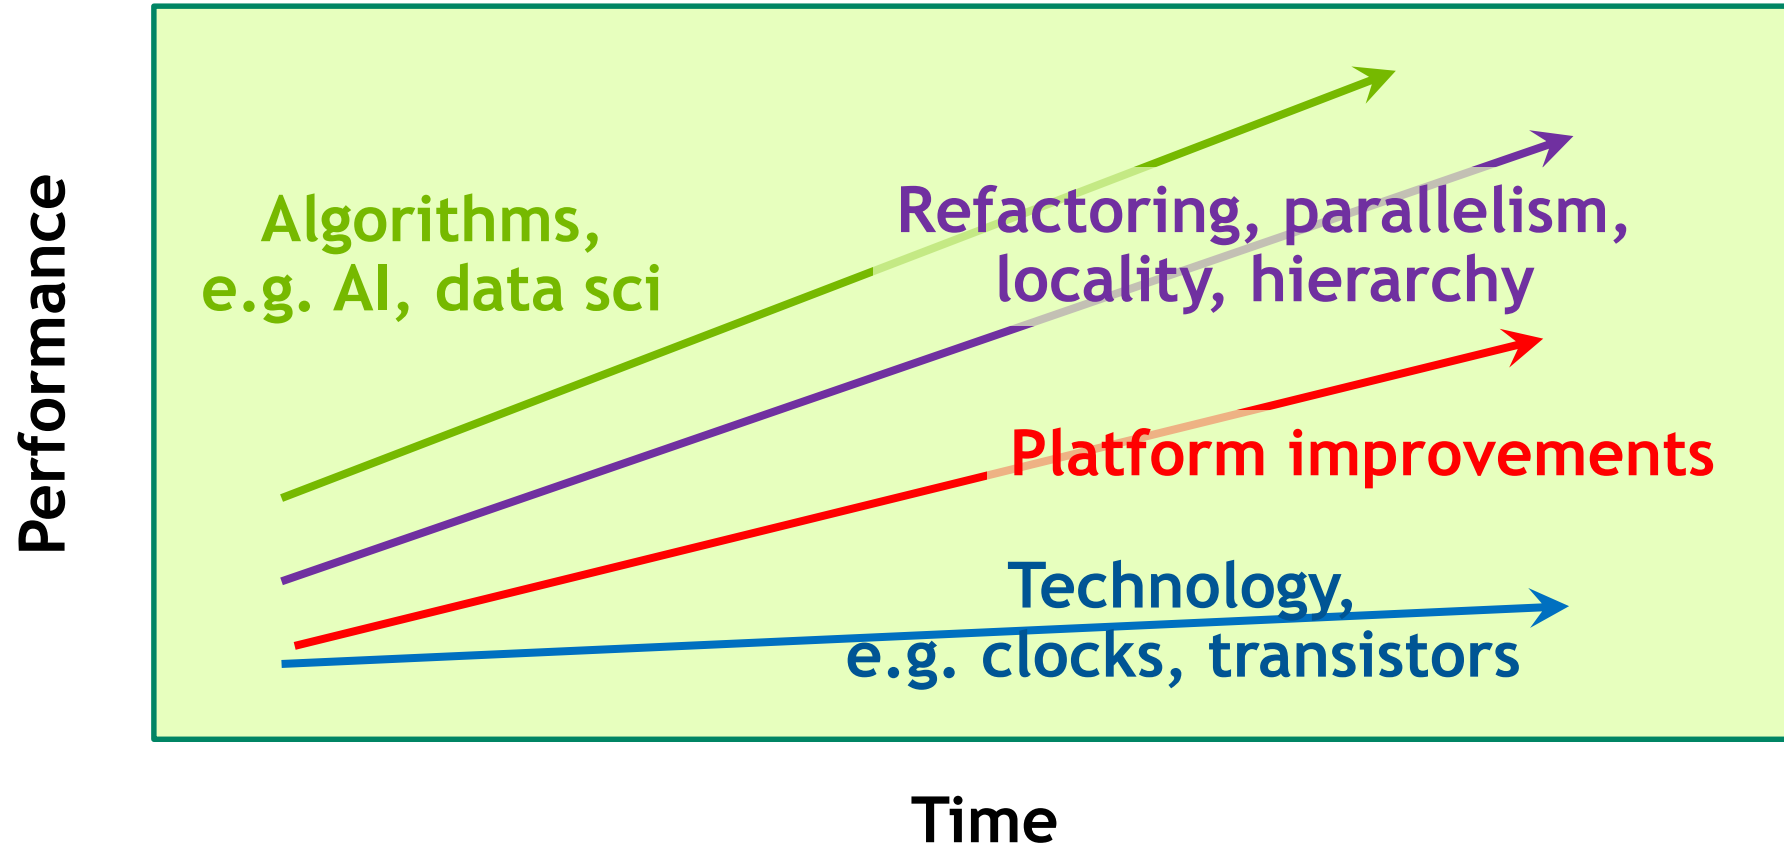
ACCELERATED COMPUTING IS FULL-STACK OPTIMIZATION

2X More Performance with Software Optimizations Alone



MATCH YOUR WORK TO A LINE


More investment, more return on that investment



CUDA PROGRAMMING ARCHITECTURE


Recent developments and directions

Improved SIMT Model



New Algorithms

Unified Memory




Languages, directives




Productive parallel programming

Tensor Core




Multi-Precision

Cooperative Groups



Hierarchical Efficiency

CUDA Graphs




Small Task Efficiency

CUDA PROGRAMMING ARCHITECTURE


Recent developments and directions

Improved SIMT Model



New Algorithms

Unified Memory




Languages, directives




Productive parallel programming

Tensor Core




Multi-Precision

Cooperative Groups



Hierarchical Efficiency


CUDA Graphs



Small Task Efficiency

THREAD CAPABILITIES

History



Thread Capabilities

Tesla: SIMT, Load/Store Pointers

Fermi: Double Precision, Function Pointers
64-bit addresses


Kepler: Dynamic Parallelism

Pascal: Inter-Block Synchronization

Volta: Independent Thread Scheduling,
Robust Memory Model


INDEPENDENT THREAD SCHEDULING

Communicating Algorithms



Pascal: Lock-Free Algorithms

Threads cannot wait for messages



Volta/Turing: Starvation Free Algorithms

Threads **may wait** for messages

[PEOPLE THINK] GPUS ARE GOOD FOR:

1. Floats, short floats, and doubles.
2. Arrays (may be multi-dimensional).
3. Coalesced memory access.
4. Lock-free algorithms.


[PEOPLE THINK] GPUS ARE BAD FOR:

1. Strings.
2. Node-based data structures.
3. Random memory walks.
4. Starvation-free algorithms (spinlocks).

CUDA PROGRAMMING ARCHITECTURE


Recent developments and directions

Improved SIMT Model



New Algorithms

Unified Memory




Languages, directives




Productive parallel programming

Tensor Core




Multi-Precision

Cooperative Groups



Hierarchical Efficiency

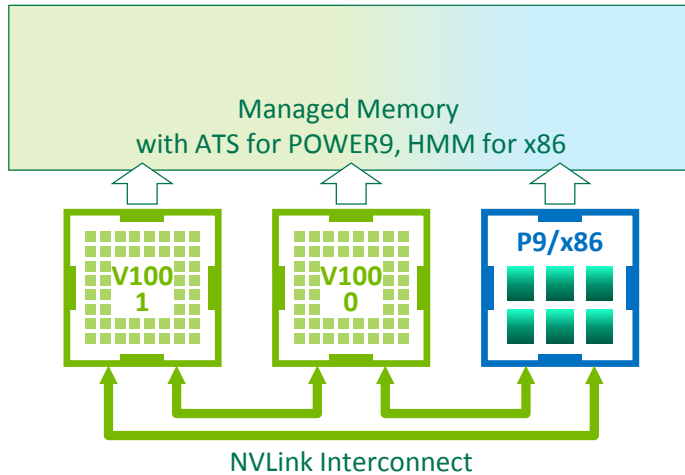
CUDA Graphs



Small Task Efficiency

UNIFIED MEMORY

Access all memory in the node



ALLOCATION

Automatic access to *all* system memory: malloc, statics, globals, stack, file system

ACCESS

All data accessible concurrently from any processor, anytime

Atomic operations resolved directly over NVLink

WHAT YOU CAN DO WITH UNIFIED MEMORY

Works everywhere today

```
int *data;  
cudaMallocManaged(&data, sizeof(int) * n);  
kernel<<< grid, block >>>(data);
```

Works on POWER9 + CUDA 9.2 and forthcoming x86 + HMM

```
int *data = (int*)malloc(sizeof(int) * n);  
kernel<<< grid, block >>>(data);
```

```
int data[1024];  
kernel<<< grid, block >>>(data);
```


```
int *data = (int*)alloca(sizeof(int) * n);  
kernel<<< grid, block >>>(data);
```

```
extern int *data;  
kernel<<< grid, block >>>(data);
```


16 GPUs WITH 32GB MEMORY EACH



16x 32GB Independent Memory Regions




NVSWITCH PROVIDES

All-to-all high-bandwidth
peer mapping between GPUs

Full inter-GPU memory
interconnect (incl. Atomics)

UNIFIED MEMORY + DGX-2



UNIFIED MEMORY PROVIDES

Single memory view
shared by all GPUs


Automatic migration of data
between GPUs

User control of data locality

CUDA PROGRAMMING ARCHITECTURE


Recent developments and directions

Improved SIMT Model



New Algorithms

Unified Memory




Languages, directives




Productive parallel programming

Tensor Core




Multi-Precision

Cooperative Groups



Hierarchical Efficiency

CUDA Graphs



Small Task Efficiency

Parallel Features in Fortran and C++

Fortran 2018

Array syntax (F90)

FORALL (F95)

Co-arrays (Fo8, F18)

DO CONCURRENT (Fo8, F18)

C++17

Threads (C++11)

pSTL Parallel Algorithms (C++17)

Fortran 2018 DO CONCURRENT Construct

Adds support for explicit shared vs private data

ISO/IEC FDIS 1539-1:2018 (E)

NOTE 5

The following code demonstrates the use of the `LOCAL` clause so that the `X` inside the `DO CONCURRENT construct` is a temporary variable, and will not affect the `X` outside the construct.

```
X = 1.0
DO CONCURRENT (I=1:10) LOCAL (X)
  IF (A (I) > 0) THEN
    X = SQRT (A (I))
    A (I) = A (I) - X**2
  END IF
  B (I) = B (I) - A (I)
END DO
PRINT *, X                                ! Always prints 1.0.
```

C++17 Parallel Algorithms (pSTL)

Parallel and Vector Concurrency for the Standard Algorithms


- Execution policies can now be applied to most STL algorithms
 - `std::execution::seq` // sequential
 - `std::execution::par` // parallel
 - `std::execution::par_unseq` // parallel+vector
- A few new parallel algorithms introduced in C++17
 - `std::transform_reduce` (POLICY, first1, last1, first2, trans_op, init, reduce_op)
 - `std::for_each_n` (POLICY, first, size, func)
 - ...

C++17 Parallel Algorithms and Beyond, Bryce Lebach, CppCon 2016

The LULESH Hydrodynamics Mini-app

LLNL Hydrodynamics Proxy (Mini-App)

- ~9000 lines of C++
- Parallel versions in MPI, OpenMP, OpenACC, CUDA, RAJA
- Designed to stress compiler vectorization, parallel overheads, on-node parallelism



codesign.llnl.gov/lulesh

C++ with OpenACC Directives

```
% pgc++ -fast -ta=tesla:managed -Minfo -c lulesh.cc
...
CalcVelocityForNodes(Domain &, double, double, int):
    2103, Accelerator kernel generated
        Generating Tesla code
    2106, #pragma acc loop gang, vector(128) ! BlockIdx.x
                                           ! threadIdx.x
...
```

```
2100 static inline
2101 void CalcVelocityForNodes(Domain &domain, const Real_t dt,
2102                           const Real_t u_cut, Index_t numNode)
2103 {
2104
2105     #pragma acc parallel loop
2106     for ( Index_t i = 0 ; i < numNode ; ++i )
2107     {
2108         Real_t xdtmp, ydtmp, zdtmp ;
2109
2110         xdtmp = domain.xd(i) + domain.xdd(i) * dt ;
2111         if( FABS(xdtmp) < u_cut ) xdtmp = Real_t(0.0);
2112         domain.xd(i) = xdtmp ;
2113
2114         ydtmp = domain.yd(i) + domain.ydd(i) * dt ;
2115         if( FABS(ydtmp) < u_cut ) ydtmp = Real_t(0.0);
2116         domain.yd(i) = ydtmp ;
2117
2118         zdtmp = domain.zd(i) + domain.zdd(i) * dt ;
2119         if( FABS(zdtmp) < u_cut ) zdtmp = Real_t(0.0);
2120         domain.zd(i) = zdtmp ;
2121     }
2122 }
```


C++17 Parallel Algorithm: `for_each_n()`

```
% pgc++ -fast -ta=tesla:managed,pstl -Minfo -c lulesh.cc  
...
```

```
2100 static inline  
2101 void CalcVelocityForNodes(Domain &domain, const Real_t dt,  
2102                             const Real_t u_cut, Index_t numNode)  
2103 {  
2104  
2105     std::for_each_n (POLICY_PAR, idx(0), numNode, [&](Index_t i)  
2106     {  
2107         Real_t xdtmp, ydtmp, zdtmp ;  
2108  
2109         xdtmp = domain.xd(i) + domain.xdd(i) * dt ;  
2110         if( FABS(xdtmp) < u_cut ) xdtmp = Real_t(0.0);  
2111         domain.xd(i) = xdtmp ;  
2112  
2113         ydtmp = domain.yd(i) + domain.ydd(i) * dt ;  
2114         if( FABS(ydtmp) < u_cut ) ydtmp = Real_t(0.0);  
2115         domain.yd(i) = ydtmp ;  
2116  
2117         zdtmp = domain.zd(i) + domain.zdd(i) * dt ;  
2118         if( FABS(zdtmp) < u_cut ) zdtmp = Real_t(0.0);  
2119         domain.zd(i) = zdtmp ;  
2120     });  
2121 }  
2122 }
```

PRODUCTIVE PROGRAMMING

Conjugate Gradient implemented using NumPy

```
import numpy as np

def solve(A, b):
    x = np.zeros(A.shape[1])
    r = b - A.dot(x)
    p = r
    rsold = r.dot(r)
    for i in range(b.shape[0]):
        Ap = A.dot(p)
        alpha = rsold / (p.dot(Ap))
        x = x + alpha * p
        r = r - alpha * Ap
        rsnew = r.dot(r)
        if np.sqrt(rsnew) < 1e-10:
            break;
        beta = rsnew / rsold
        p = r + beta * p
        rsold = rsnew
    return x
```

```
import legate as np

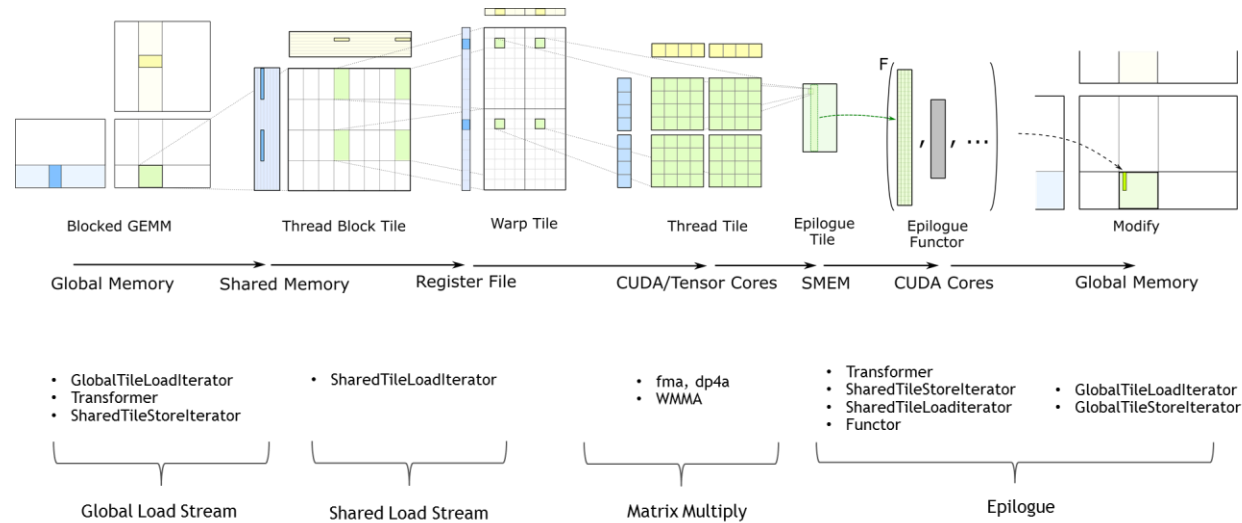
def solve(A, b):
    x = np.zeros(A.shape[1])
    r = b - A.dot(x)
    p = r
    rsold = r.dot(r)
    for i in range(b.shape[0]):
        Ap = A.dot(p)
        alpha = rsold / (p.dot(Ap))
        x = x + alpha * p
        r = r - alpha * Ap
        rsnew = r.dot(r)
        if np.sqrt(rsnew) < 1e-10:
            break;
        beta = rsnew / rsold
        p = r + beta * p
        rsold = rsnew
    return x
```

On CPU or GPU system: \$./legate app.py

C++ TEMPLATES: CUTLASS 1.1

High-performance Matrix Multiplication in Open Source CUDA C++

- ▶ Turing optimized GEMMs
 - ▶ Integer (8-bit, 4-bit and 1-bit) using WMMA
- ▶ Batched strided GEMM
- ▶ Support for CUDA 10.0
- ▶ Updates to documentation and more examples




CUTLASS GEMM Structural Model

CUTLASS 1.1

High-performance Matrix Multiplication in Open Source CUDA C++

- ▶ Turing optimized GEMMs
 - ▶ Integer (8-bit, 4-bit and 1-bit) using WMMA
- ▶ Batched strided GEMM
- ▶ Support for CUDA 10.0
- ▶ Updates to documentation and more examples


<https://github.com/NVIDIA/cutlass>



CUDA PROGRAMMING ARCHITECTURE


Recent developments and directions

Improved SIMT Model



New Algorithms

Unified Memory




Languages, directives




Productive parallel programming

Tensor Core




Multi-Precision

Cooperative Groups



Hierarchical Efficiency

CUDA Graphs



Small Task Efficiency


REDUCED PRECISION IN HPC

Why?


1. Reduce memory traffic
2. Reduce memory footprint
3. Utilize accelerated hardware

....without compromising the end result


PRECISION MATTERS FOR HPC AND AI



VOLTA TENSOR CORE GPU



MULTI-PRECISION
COMPUTING



FUSION OF HPC & AI

REDUCED PRECISION IN HPC

Mixed-precision algorithms are ever more popular

Most common use case has been combining double + single precision...
...but not the only choice

Combining floating point and integer operations for performance and determinism

- E.g., AMBER - fp32 pair interactions with int64 energy summation

Use fp16 or int16 precision to reduce memory traffic

Use accelerated fp16 or tensor-core instructions for *large* speedups

- E.g., tensor-core accelerated mixed-precision LU solver (Dongarra *et al*)


SPARSE SOLVERS AND LOW PRECISION

Since sparse solvers are iterative, judicious use of precision can lead to large speedups

Most sparse solvers are memory bandwidth bound, so reduced precision accelerates through reduced memory traffic

Recipes need to be worked out

- Wrong algorithm will not converge
- E.g., Defect correction vs reliable updates
- Keep the precision where it's needed




MILC/QUDA CG solver
Condition number $\sim 10^6$

EXTREME EXAMPLE

Chroma/QUDA multigrid solver uses 5 different precisions

- double - outer defect correction
- single - GCR solver
- half - preconditioner
- quarter - preconditioner halo communication
- int32 - deterministic parallel coarsening

Reduced precision critical for memory traffic and footprint reduction




Chroma HMC benchmark (Bálint Joó, Jlab)

Machine x Algorithm = 73x speedup

ECP >50x goal achieved pre-Exascale

CUDA TENSOR CORE PROGRAMMING

16x16x16 Warp Matrix Multiply and Accumulate (WMMA)




$$\mathbf{D} = \begin{pmatrix} A_{0,0} & A_{0,1} & A_{0,\dots} & A_{0,15} \\ A_{1,0} & A_{1,1} & A_{1,\dots} & A_{1,15} \\ A_{\dots,0} & A_{\dots,1} & A_{\dots,\dots} & A_{\dots,15} \\ A_{15,0} & A_{15,1} & A_{15,\dots} & A_{15,15} \end{pmatrix} \begin{pmatrix} B_{0,0} & B_{0,1} & B_{0,\dots} & B_{0,15} \\ B_{1,0} & B_{1,1} & B_{1,\dots} & B_{1,15} \\ B_{\dots,0} & B_{\dots,1} & B_{\dots,\dots} & B_{\dots,15} \\ B_{15,0} & B_{15,1} & B_{15,\dots} & B_{15,15} \end{pmatrix} + \begin{pmatrix} C_{0,0} & C_{0,1} & C_{0,\dots} & C_{0,15} \\ C_{1,0} & C_{1,1} & C_{1,\dots} & C_{1,15} \\ C_{\dots,0} & C_{\dots,1} & C_{\dots,\dots} & C_{\dots,15} \\ C_{15,0} & C_{15,1} & C_{15,\dots} & C_{15,15} \end{pmatrix}$$

FP16 or FP32 FP16 FP16 FP16 or FP32

$$\mathbf{D} = \mathbf{AB} + \mathbf{C}$$

LINEAR ALGEBRA + TENSOR CORES



Double Precision LU Decomposition

- Compute initial solution in FP16
- Iteratively refine to obtain same accuracy as FP64

Achieved FP16->FP64 Tflops: **24**

Achieved FP64 Tflops: **5.8**

Device FP64 Tflops: **7.5**

MAGMA Library with FP16


Data courtesy of: Azzam Haidar, Stan. Tomov & Jack Dongarra, Innovative Computing Laboratory, University of Tennessee
“Harnessing GPU Tensor Cores for Fast FP16 Arithmetic to Speed up Mixed-Precision Iterative Refinement Solvers”,
A. Haidar, S. Tomov, J. Dongarra, N. Higham SC’18
GTC 2018 Poster P8237: Harnessing GPU’s Tensor Cores Fast FP16 Arithmetic to Speedup Mixed-Precision Iterative Refinement Solves



ADVANCING FUSION DISCOVERIES

ASGarD: Adaptive Sparse Grid Discretization

Two stream instability study



Scientists believe fusion is the future of energy but maintaining plasma reactions is challenging and disruptions can result in damage to the tokamak. Researchers at ORNL are simulating instabilities in the plasma to provide physicists a better understanding of what happens inside the reactor.

With NVIDIA Tensor Cores the simulations run 3.5X faster

than previous methods so the team can simulate significantly longer physical times and help advance our understanding of how to sustain the plasma and generate energy




Joint work with ORNL & UTK David Green, Ed Azevedo, Wael Elwasif, Graham Lopez, Tyler McDaniel, Lin Mu, Stan Tomov, Jack Dongarra

CUDA PROGRAMMING ARCHITECTURE


Recent developments and directions

Improved SIMT Model



New Algorithms

Unified Memory




Languages, directives




Productive parallel programming

Tensor Core




Multi-Precision

Cooperative Groups



Hierarchical Efficiency

CUDA Graphs



Small Task Efficiency

CUDA 9+: COOPERATIVE GROUPS

Cooperative Group

```
graph TD; A[Cooperative Group] --> B[The ability to synchronize and coordinate across explicit granularities in the program.]; A --> C[Program-defined values (i.e. opaque handles) that represent a set of threads.]; B --> D[Available horizontal operations across its named set of threads (e.g. collectives).]; C --> D;
```

The ability to synchronize and coordinate across explicit granularities in the program.

Program-defined values (i.e. opaque handles) that represent a set of threads.

Available horizontal operations across its named set of threads (e.g. collectives).

SCOPE OF COOPERATION

Warp

For currently converged threads:
`auto group = cg::coalesced_threads();`

SM

For CUDA thread blocks:
`auto group = cg::this_thread_block();`

GPU

For cooperative grids:
`auto group = cg::this_grid();`

Node

For cooperative grids that span devices:
`auto group = cg::this_multi_grid();`



Cooperating on an SM



Cooperating across a GPU


Cooperating across
many GPUs with NVLINK

SYNCHRONIZE AT ANY SCALE

Three Key Capabilities


FLEXIBLE GROUPS

Define and
synchronize arbitrary
groups of threads




WHOLE-GRID SYNCHRONIZATION

Synchronize multiple
thread blocks




MULTI-GPU SYNCHRONIZATION



CUDA PROGRAMMING ARCHITECTURE


Recent developments and directions

Improved SIMT Model



New Algorithms

Unified Memory




Languages, directives




Productive parallel programming

Tensor Core




Multi-Precision

Cooperative Groups



Hierarchical Efficiency

CUDA Graphs



Small Task Efficiency

ASYNCHRONOUS GRAPHS

Execution optimization when workflow is known up front

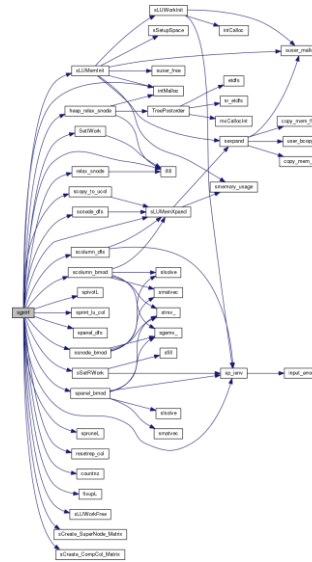
```
// Basic function to test primality.
bool IsPrime( size_t n)
{
    if (n == 2) return true;
    if ((n == 1) || ((n % 2) == 0)) return false;
    size_t iters = (unsigned int)sqrt((double)n);
    for (size_t i = 3; i <= iters; i+=2) if (n % i == 0) return false;
    return true;
}

// Compute primes from 1 to 100,000,000.
size_t ComputePrimes()
{
    size_t Primes = 0;
    for ( size_t Start = 1; Start <= 100000000; ++Start)
    {
        if ( IsPrime( Start) )
        {
            ++Primes;
        }
    }
    return Primes;
}
```


Loop & Function
offload




DL Inference



Linear Algebra




Deep Neural Network
Training



HPC Simulation


THE GRAPH ADVANTAGE

WHOLE WORKFLOW OPTIMIZATIONS



Seeing all work at once enables new optimizations in hardware and software

EFFICIENT LAUNCH OF COMPLEX WORK



Launch potentially thousands of work items with a single call

DEFINITION OF A CUDA GRAPH

Graph Nodes Are Not Just Kernel Launches

Sequence of operations, connected by dependencies.

Operations are one of:

Kernel launch

CUDA kernel running on GPU

CPU function call


Callback function on CPU

Memcpy/Memset

GPU data management

Sub-Graph


Graphs are hierarchical



THREE-STAGE EXECUTION MODEL

Early structure binding amortizes overheads, late data binding enables iteration

Define




Single Graph “Template”

Mutable structure

Created in host code
or built up from libraries

Instantiate




Multiple “Executable Graphs”

Fixed structure, mutable data

Snapshot of template
Sets up & initializes GPU
execution structures
(create once, run many times)

Execute



Executable Graphs
Running in CUDA Streams
Fixed structure and data

Concurrency in graph
is **not** limited by stream

GPUDIRECT FAMILY

Technologies enabling products

GPUDIRECT SHARED GPU-SYMEM

GPU pinned memory shared with other
RDMA-capable devices
Avoids intermediate copies

GPUDIRECT P2P

Accelerated GPU-GPU memory copies
Inter-GPU direct load/store access

GPUDIRECT RDMA

Direct GPU to 3rd party device transfers
E.g. direct I/O, optimized inter-node
communication

GPUDIRECT ASYNC

Direct GPU to 3rd party device synchronizations
E.g. optimized inter-node communication

NCCL

NVIDIA Collective Communications Library

Optimized collective communication library between CUDA devices.

Easy to integrate into any DL framework, as well as traditional HPC apps using MPI.

Runs on the GPU using asynchronous CUDA kernels, for faster access to GPU memory, parallel reductions, NVLink usage.

Similar to MPI collectives, but has a CUDA stream parameter, operates on CUDA pointers.

- Enables use of multiple threads per MPI rank, multiple GPUs per thread.

- Minimizes GPU threads to permit other computation to progress simultaneously.


NCCL 2 adds support for inter-node communication using Sockets or IB verbs.

Uses staging buffers on GPU within each process, so no IPC overhead

NCCL 2.0


Provide the maximum bandwidth

Allreduce Bandwidth (OMB, size=128MB, in GB/s)



NCCL 2.1

Latency improvement



NVSHMEM

Programming abstraction for communication among GPUs

shmem_init - set up communication, including CUDA P2P, IPC

shmem_malloc - get pointer to perform allocation on local GPU


shmem_ptr - given a PE, get the base address of the symmetric heap on remote GPU

Inside or outside of GPU kernel

ld/st using offsets relative to that base address

put/get p/g for abstracted communications

collectives



CONCLUSION

THE NEW HPC WORKLOAD

A Dramatic Shift from the Past




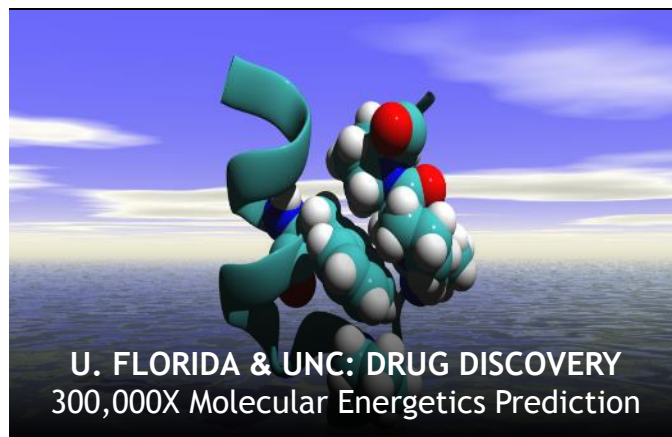

PREVIOUS NORMAL

- Simulated data sources
- Conventional 64 bit full order ab initio simulation
- Conventional reduced order methods
- All jobs are batch

NEW NORMAL

- A mix of experimental and simulated/emulated data sources
- Conventional 64 bit full order ab initio simulation
- Mixed precision algorithms with new methods for ab initio simulation
- Embedded neural nets within full order ab initio simulation
- Conventional reduced order methods
- Deep neural net and machine learning emulation
- Interactive large scale training (model parallel, data parallel and hyperparameter optimization)


CONVERGED HPC*AI EXAMPLES IN ALL DOMAINS



2X HIGHER VALUE IN 2 YEARS

Same Hardware More Performance With Software Stack Optimizations



AMBER
Chroma
GROMACS
GTC
LAMMPS
MILC
NAMD
QE
RTM
SPECFEM3D
VASP





APPLICATIONS ON PATH OF ACCELERATION

Because status-quo yields only incremental gain

Relative to dual
Haswell



Relative to dual
Haswell



GORDON BELL NOMINEES ILLUSTRATE THE NEW NORMAL


New Algorithms

Scalable HPC*AI

New Algorithms
HPC*AI


Scalable HPC*AI

New Algorithms




OAK RIDGE
National Laboratory

Genomics
2.36 ExaOps



nvidia

Weather
1.15 ExaOps



東京大学
THE UNIVERSITY OF TOKYO


OAK RIDGE
National Laboratory

Seismic
25X



OAK RIDGE
National Laboratory

Material Science
300X



Lawrence Livermore
National Laboratory

nvidia

Quantum
Chromodynamics
15x

