ENDURING DIFFERENTIATION

Timothy Lanfear

NVIDIA

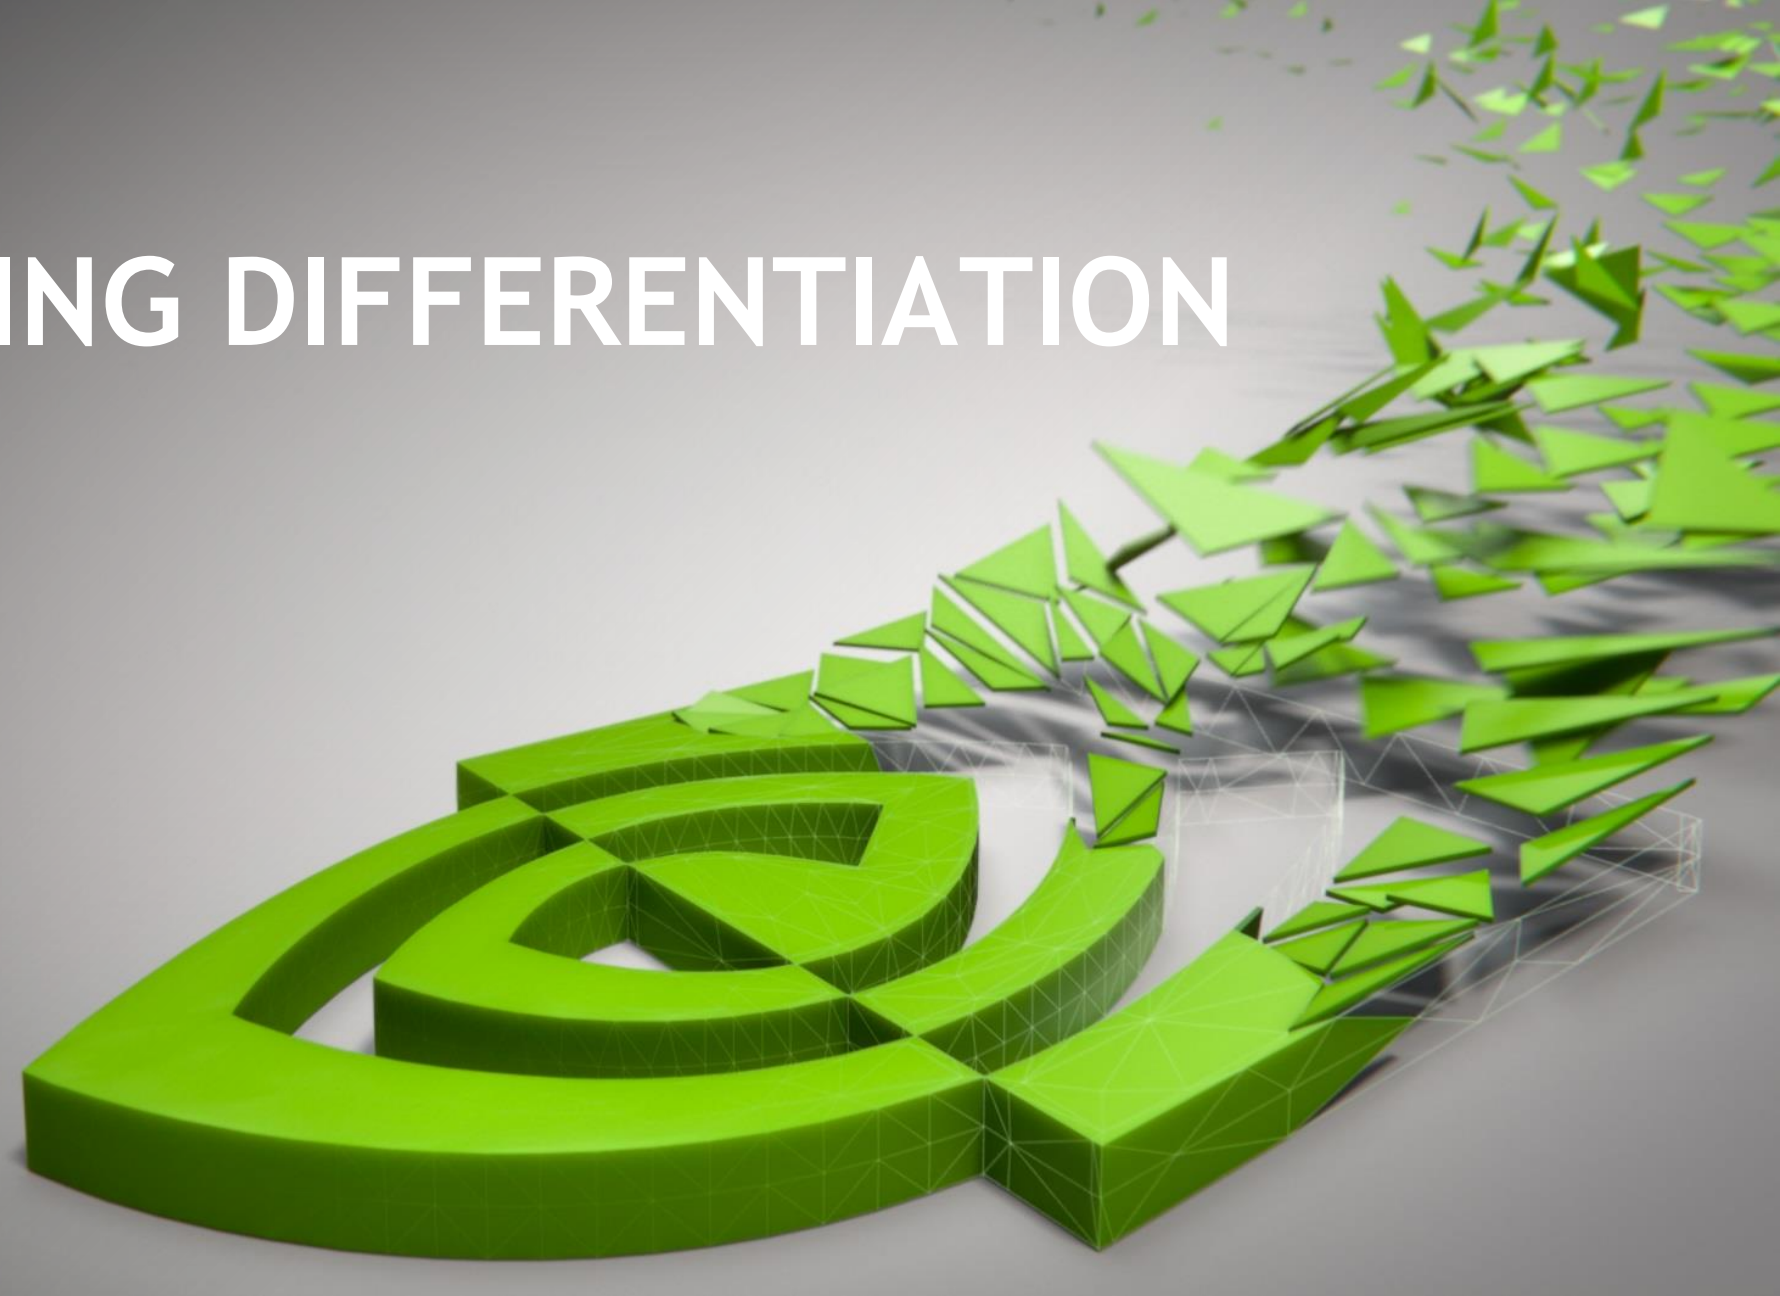

WHERE ARE WE?

LIFE AFTER DENNARD SCALING

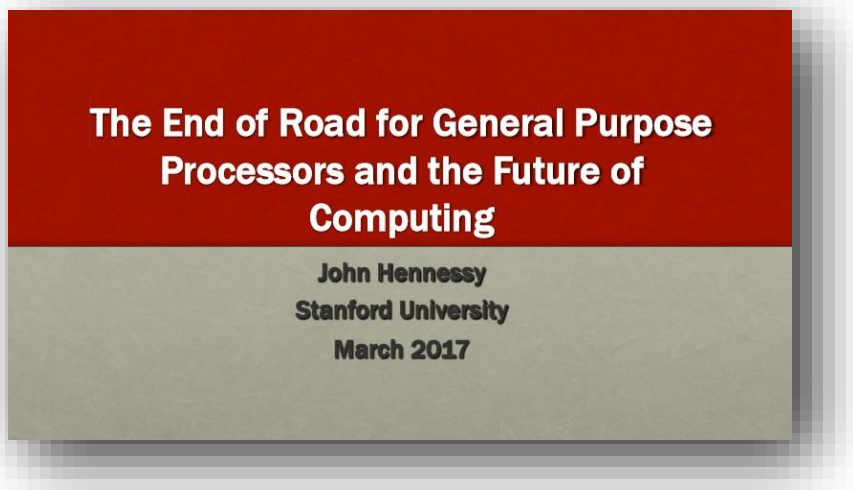

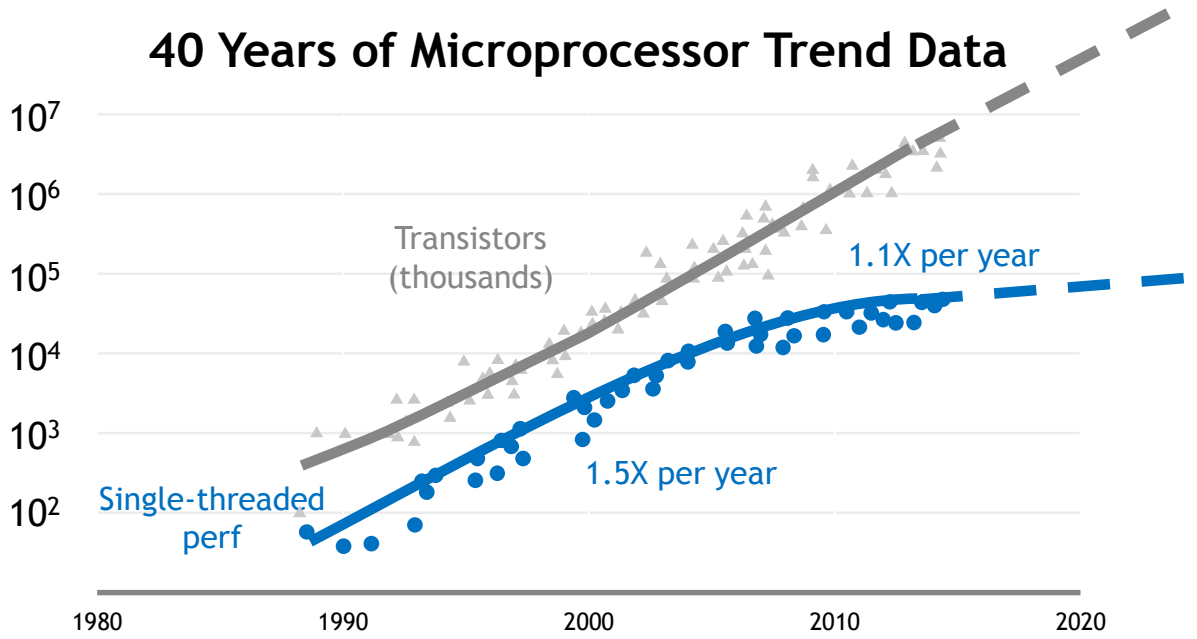

Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten New plot and data collected for 2010-2015 by K. Rupp

GPU-ACCELERATED PERFORMANCE

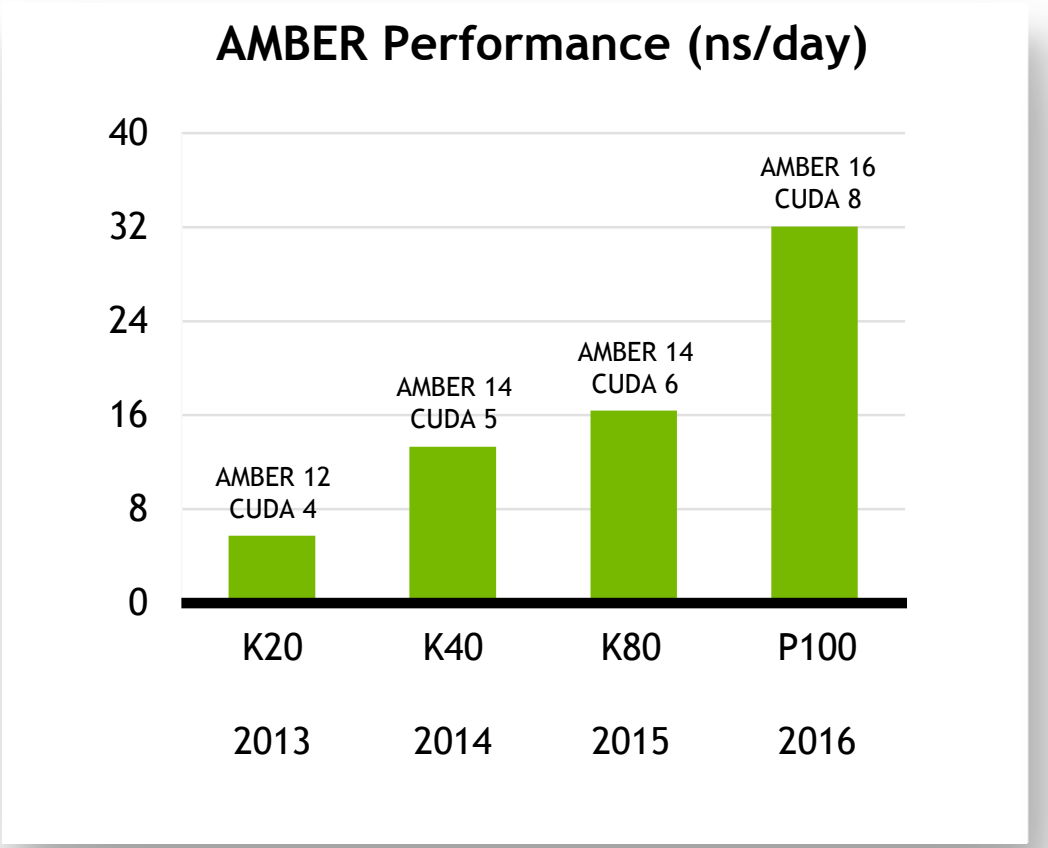

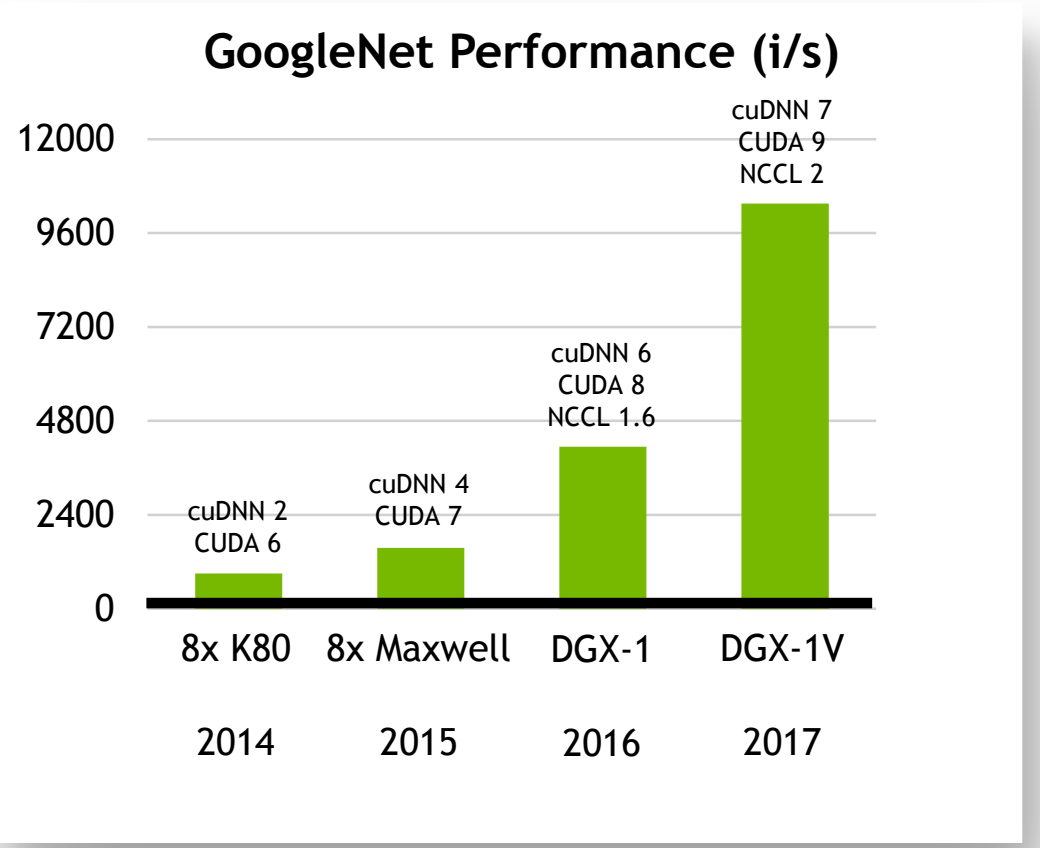

TESLA PLATFORM ADVANTAGE

Delivered value grows over time

Life Sciences (AMBER)

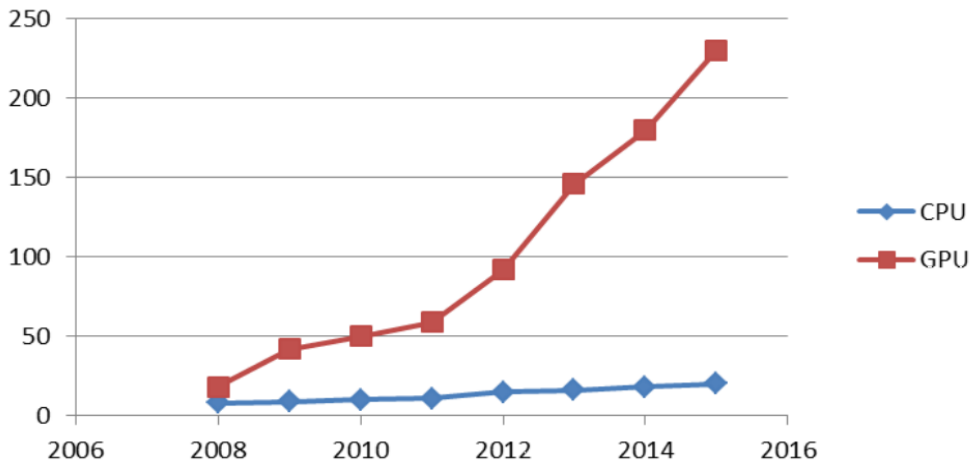

10X Perf in 8 Years

Oil & Gas (RTM)

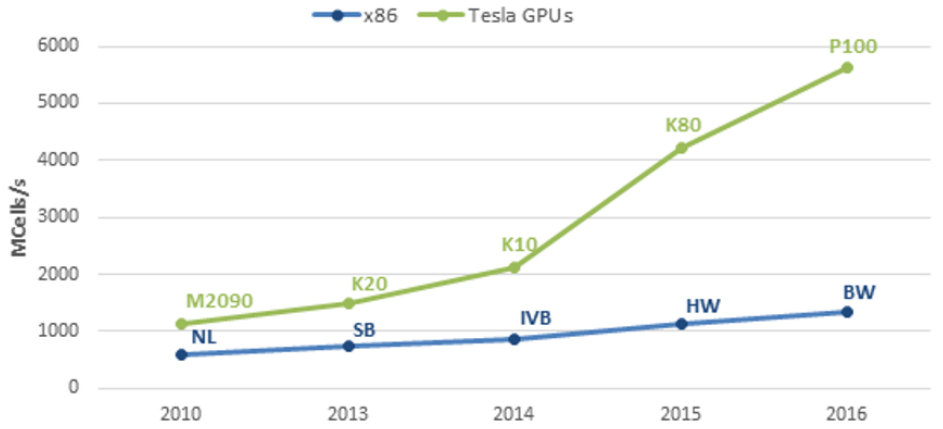

6X Perf in 6 Years

Amber performance: Nano Seconds Per Day delivered on 1xServer with GPUs and CPUS

GPU-ACCELERATED EFFICIENCY

13/13 Greenest Supercomputers
Powered by Tesla P100

TSUBAME 3.0
Kukai
AIST AI Cloud
RAIDEN GPU subsystem
Piz Daint
Wilkes-2
GOSAT-2 (RCF2)
DGX Saturn V
Reedbush-H
JADE
Facebook Cluster
Cedar
DAVIDE

On Track To Meet Exascale Goal

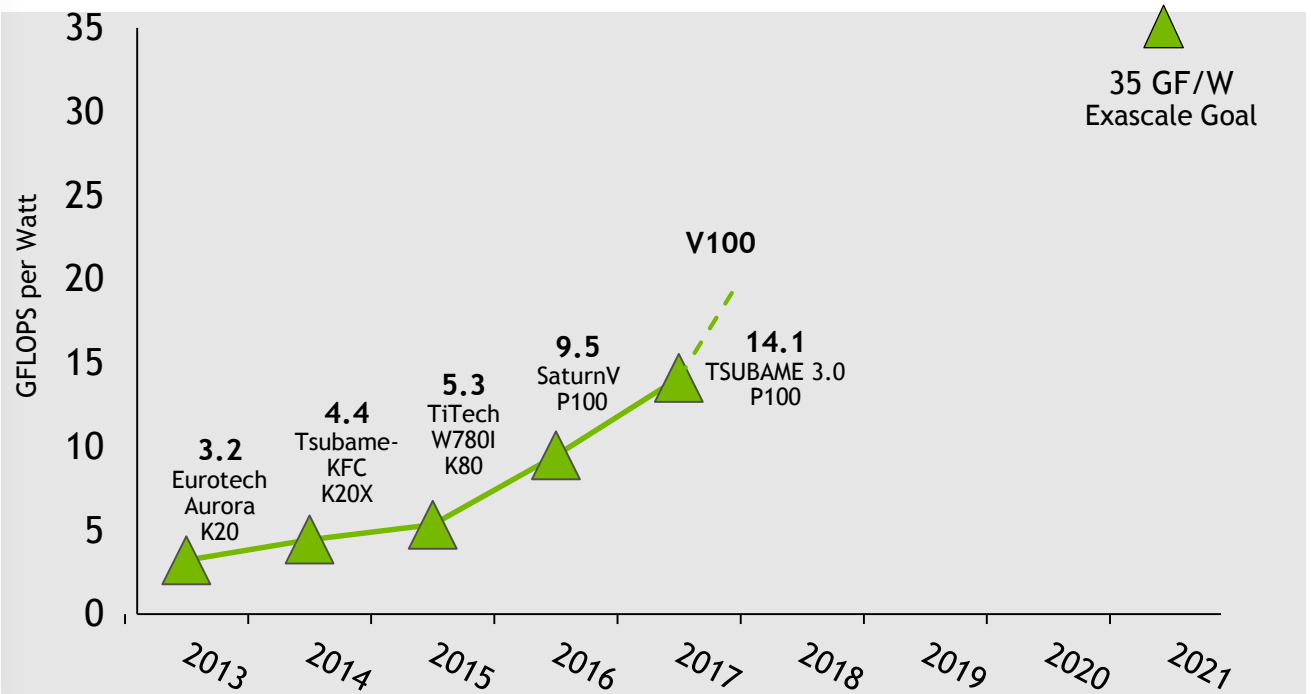

Top GPU Systems in Green500 Lists and NVIDIA Projections for V100

HOW ARE WE DOING THIS?

And, *is our differentiation sustainable*?

- What are the most important dimensions of our differentiation?
- Why are GPUs so much more efficient than CPUs?
- How can we continue scaling performance/efficiency as Moore's Law fades?
- Why can't competitors replicate GPU efficiency, performance, scaling, etc., with lots of weak CPU cores? (e.g., Intel KNC/KNL/KNM)
- How is optimizing GPUs for AI affecting their suitability for HPC?

ENERGY EFFICIENCY

COMPUTATION VERSUS COMMUNICATIONS

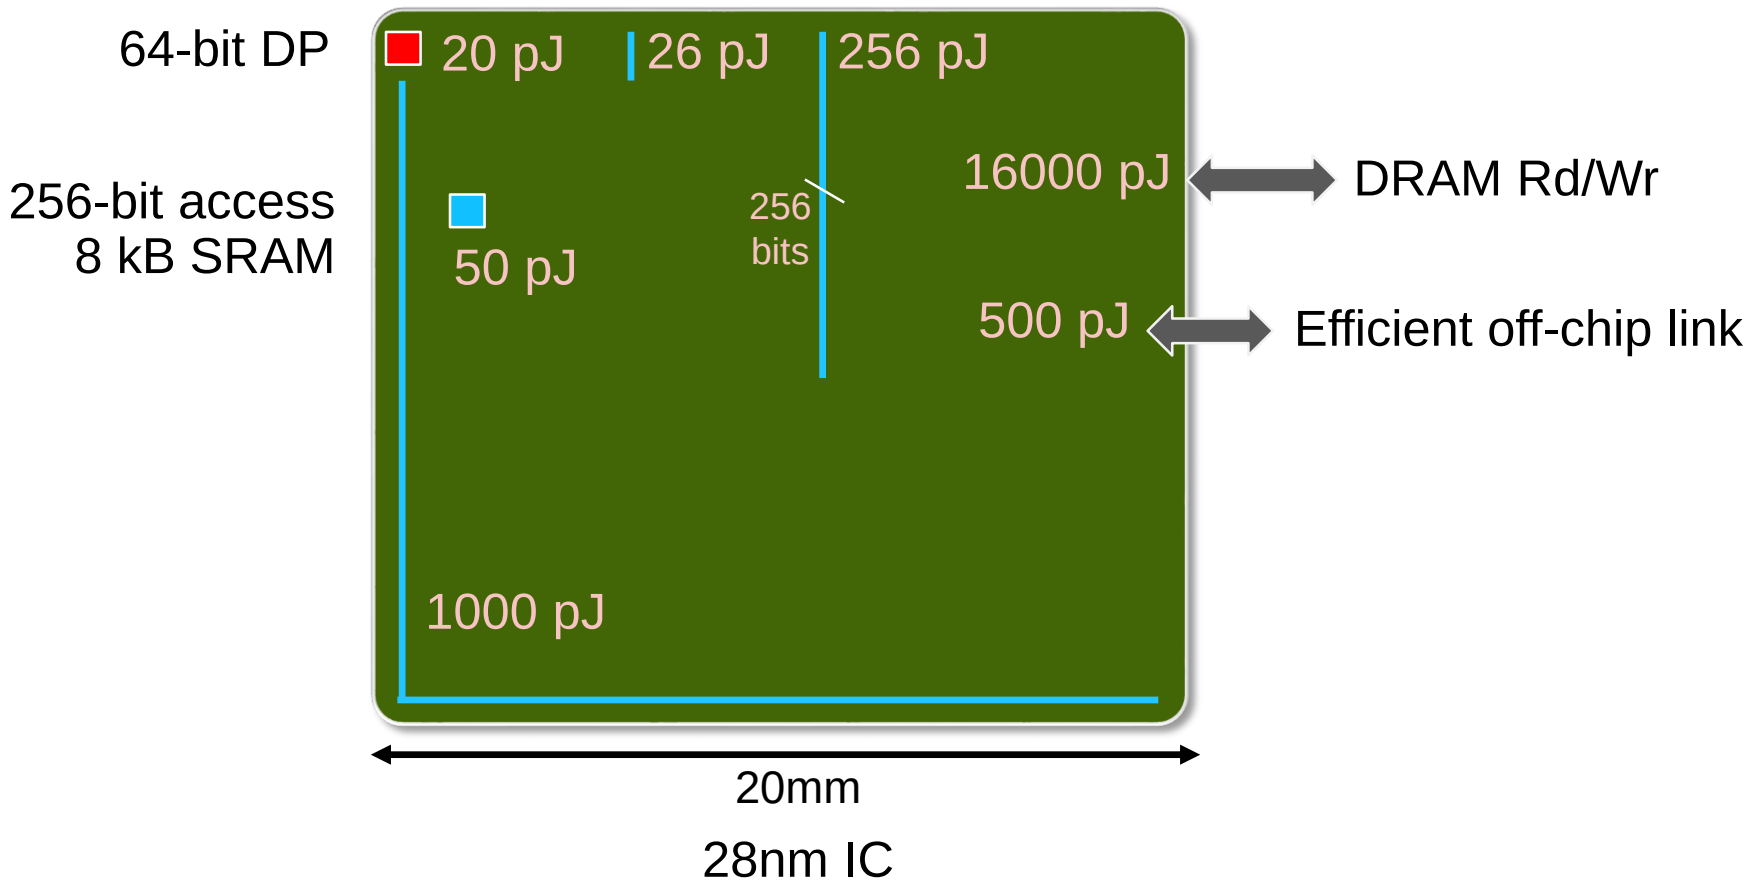

CPU
126 pJ/flop (SP)

Optimized for Latency

Deep Cache Hierarchy

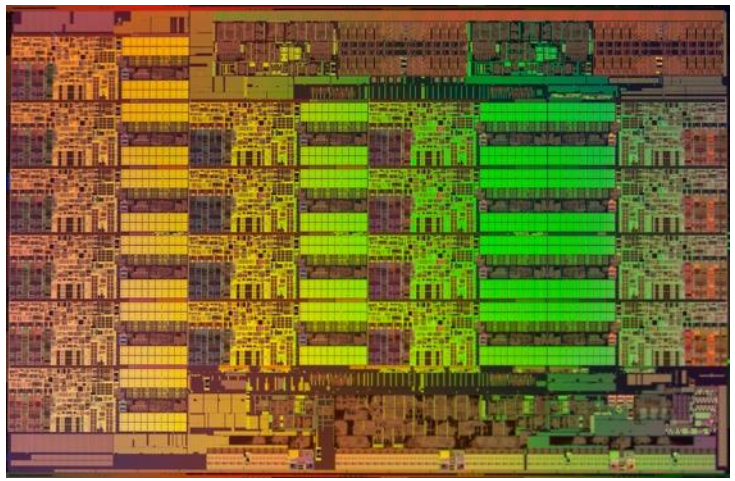

Broadwell E5 v4
14 nm

GPU
28 pJ/flop (SP)

Optimized for Throughput

Explicit Management
of On-chip Memory

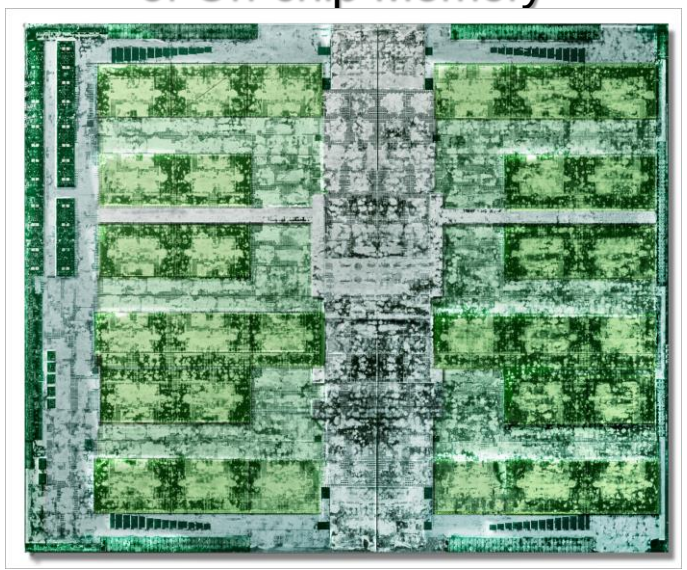

Pascal
16 nm

HOW IS POWER SPENT IN A CPU?

In Order, Embedded

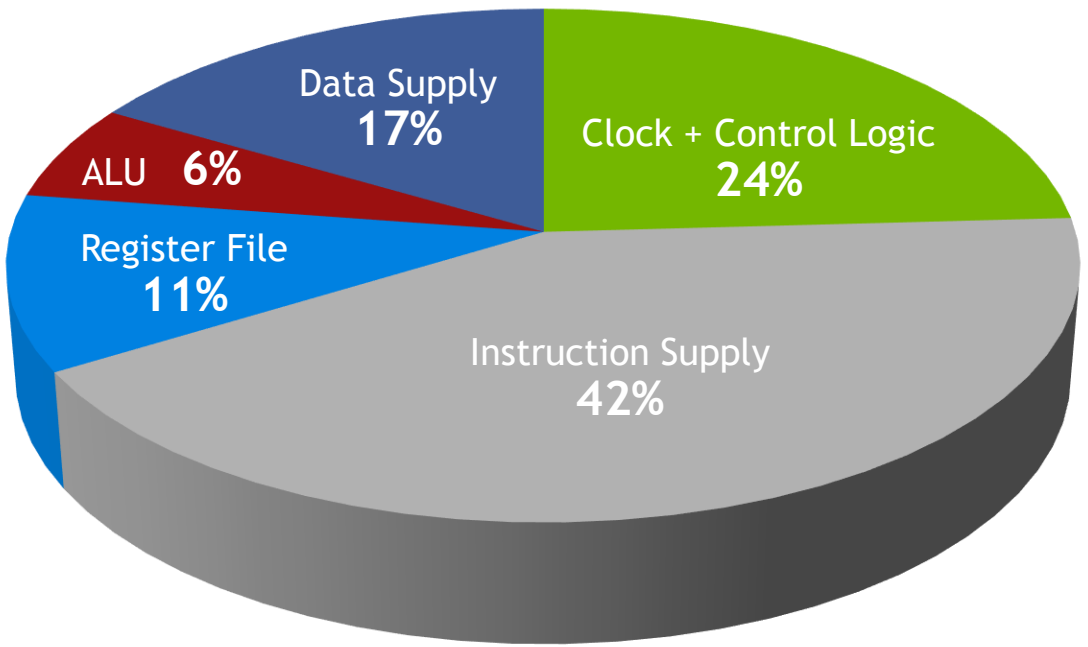

Dally [2008] (Embedded in-order CPU)

Out of Order, High Performance

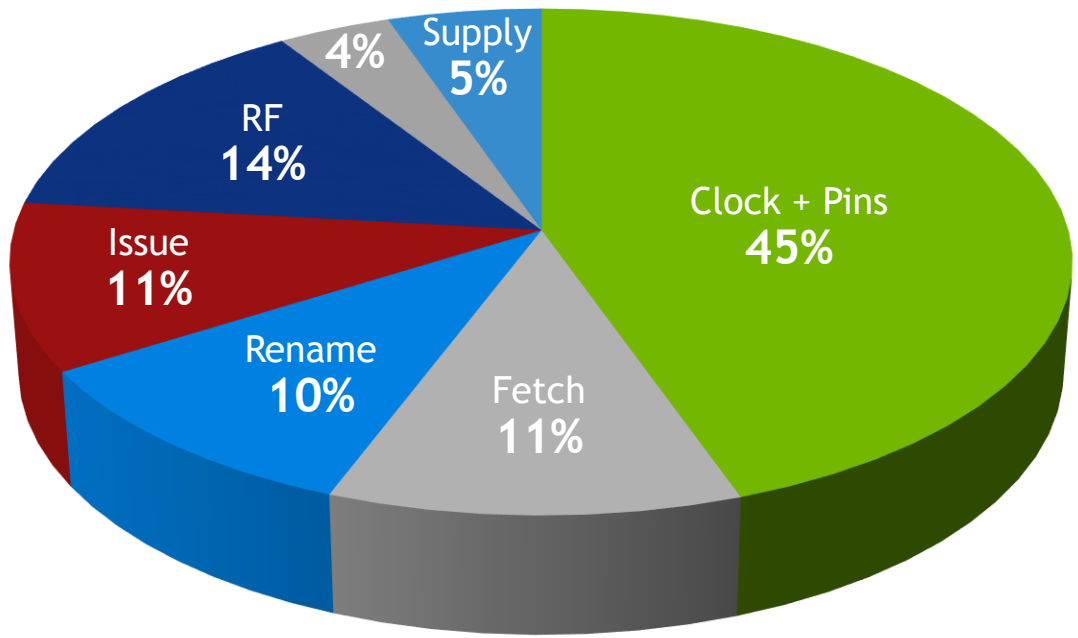

Natarajan [2003] (Alpha 21264)

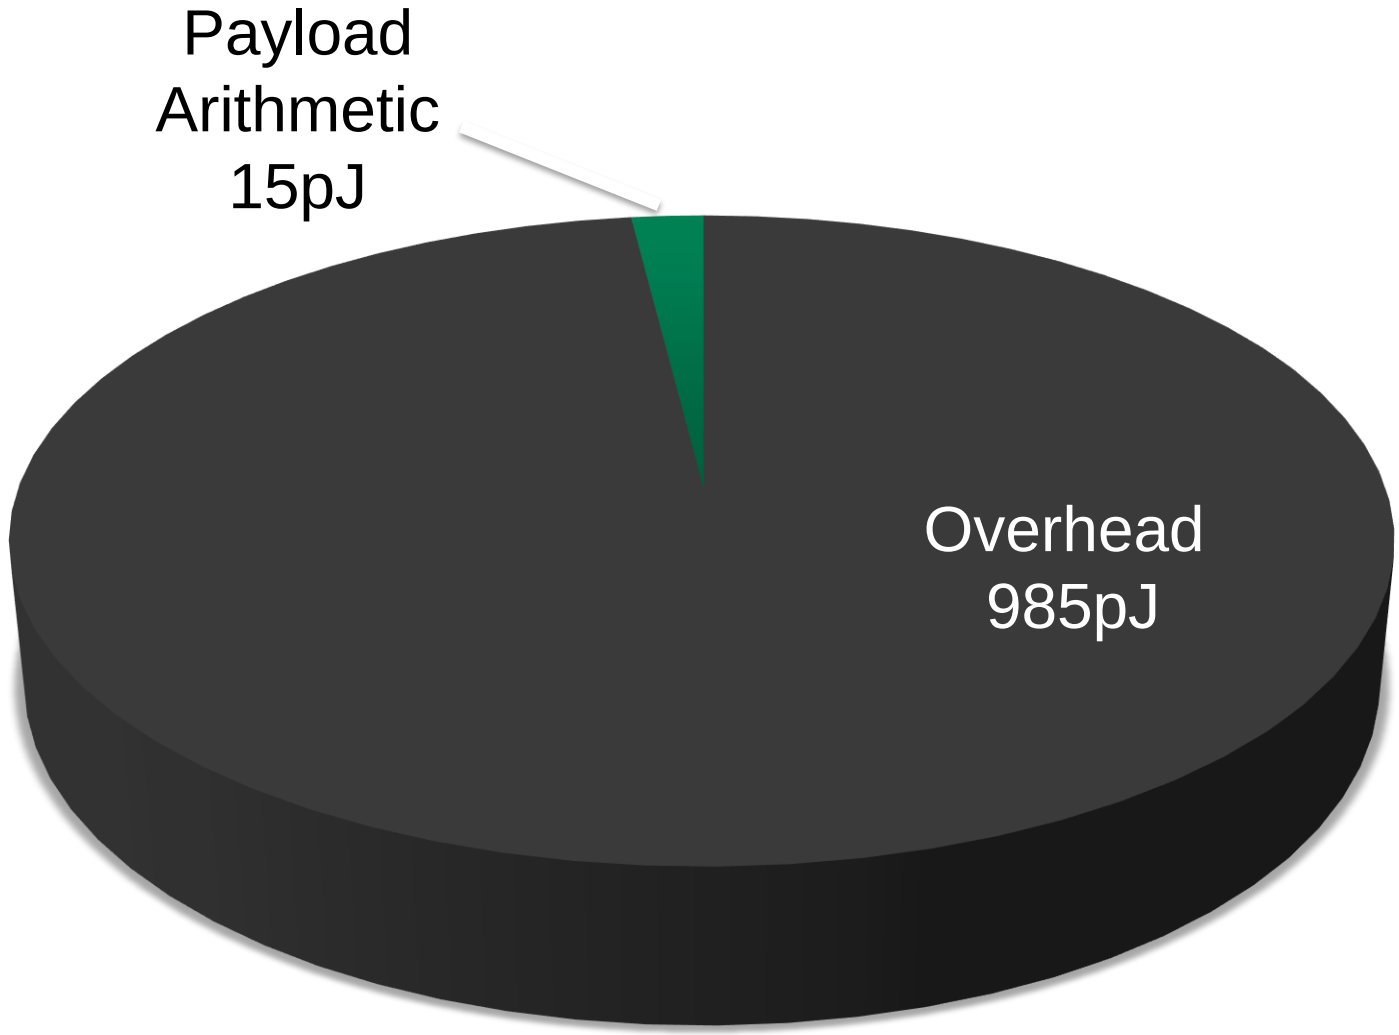
Payload
Arithmetic
15pJ
Overhead
985pJ

SIMPLER CORES = ENERGY EFFICIENCY

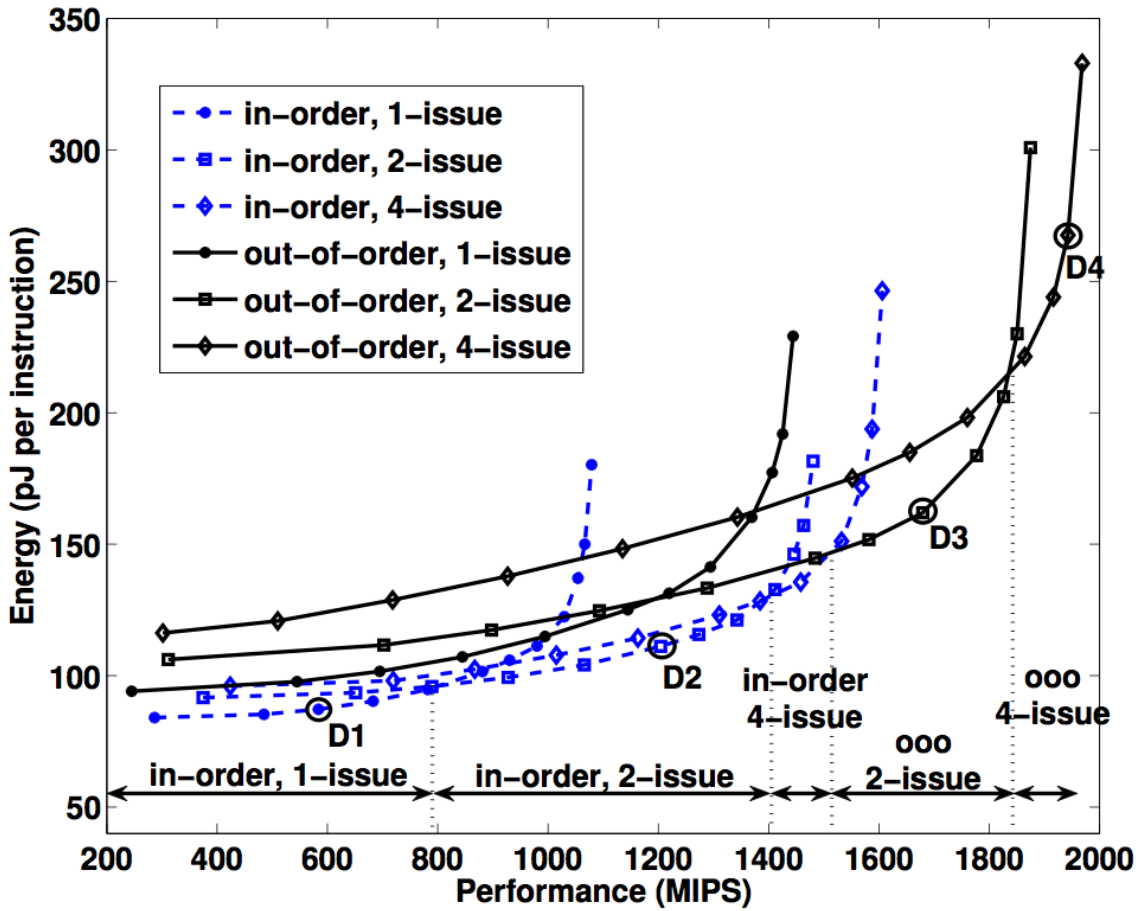

Source: Azizi [PhD 2010]

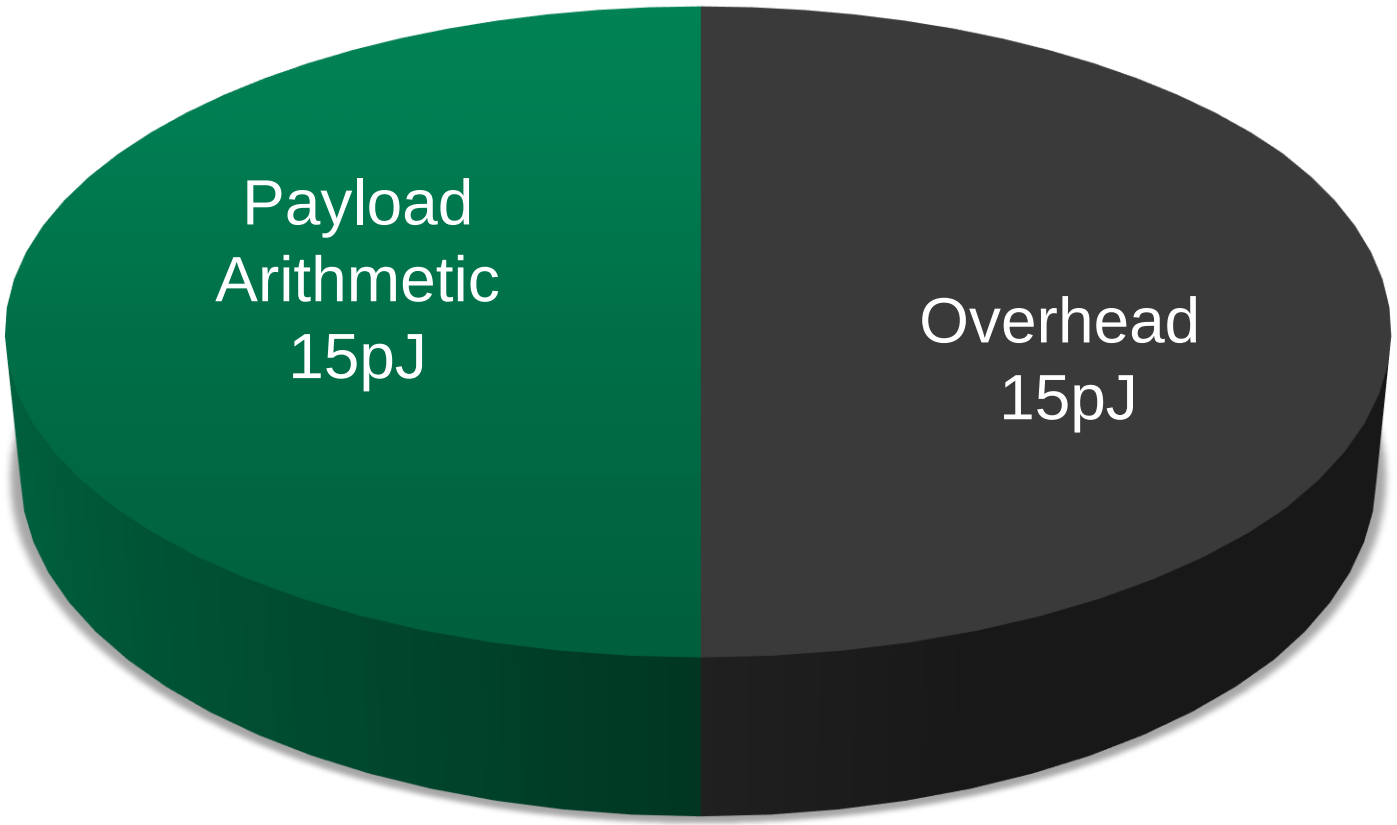
Payload
Arithmetic
15pJ
Overhead
15pJ

THROUGHPUT PROCESSORS

RISE OF LEAKAGE

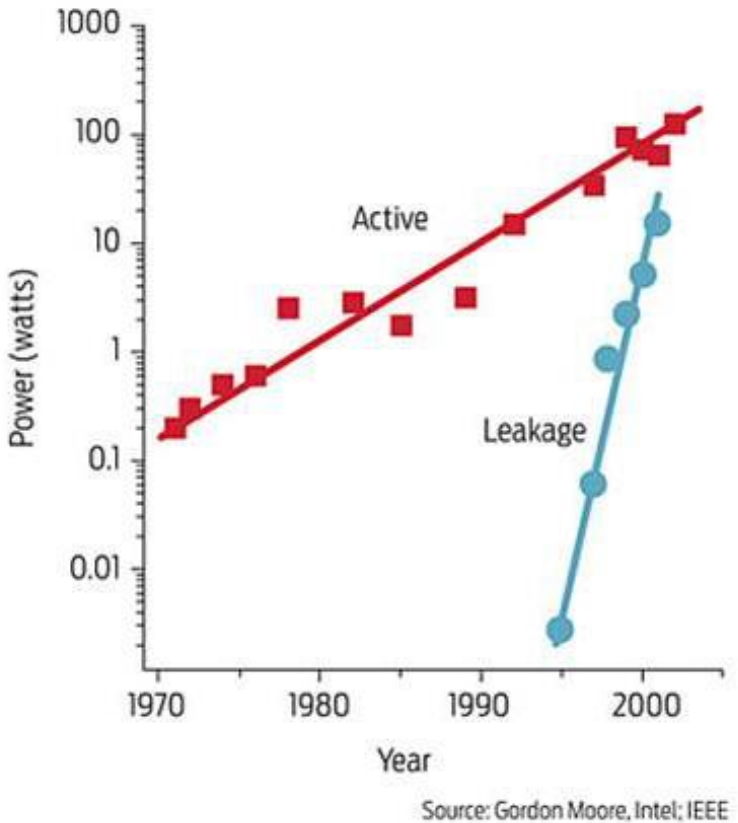

Source: Gordon Moore, Intel; IEEE

FREQUENCY VS. LEAKAGE

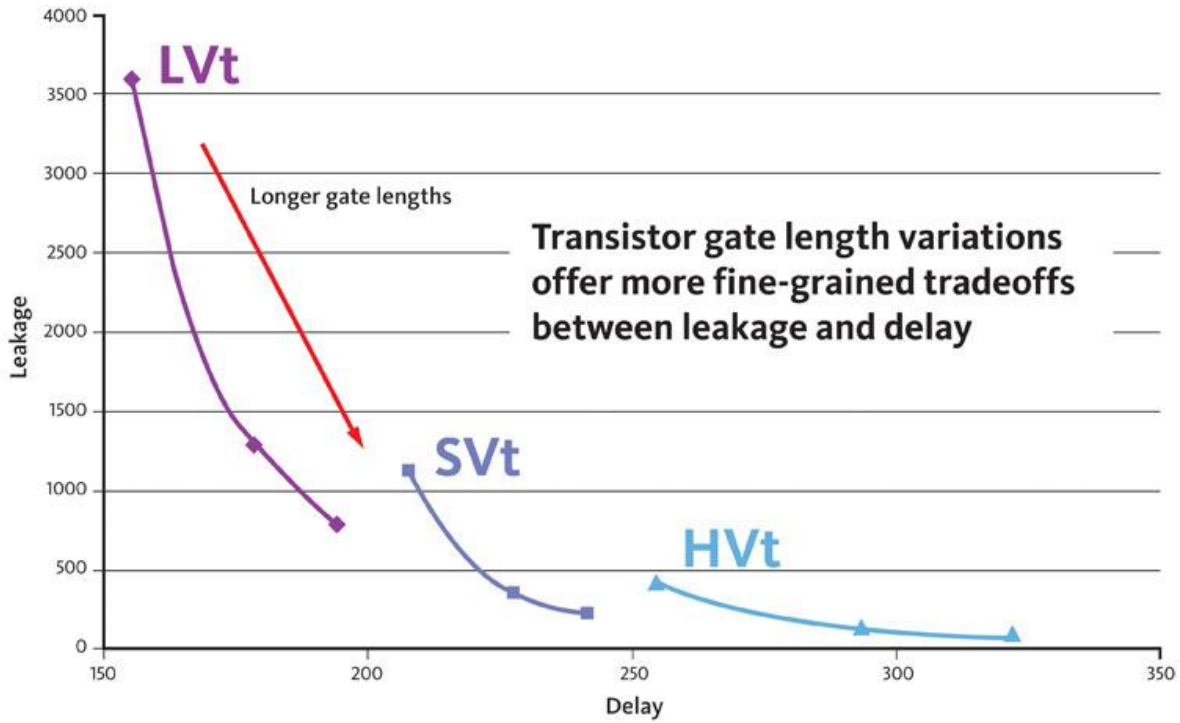

Source: Gordon Moore, Intel; IEEE

OPTIMIZED FOR DATACENTER EFFICIENCY

40% More Performance in a Rack

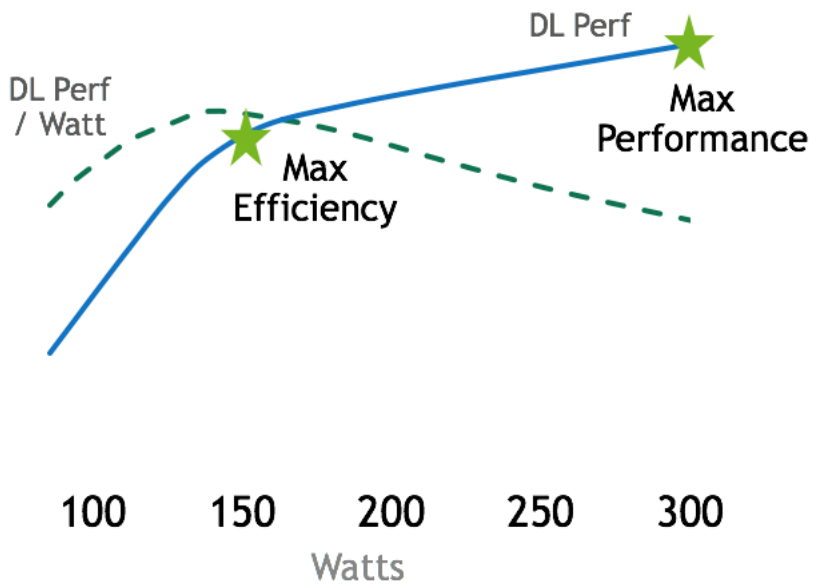

80% Perf at Half the Power

V100 Max Performance	V100 Max Efficiency
13 KW Rack 4 Nodes of 8xV100	13 KW Rack 7 Nodes of 8xV100
13 ResNet-50 Networks Trained Per Day	18 ResNet-50 Networks Trained Per Day

ResNet-50 Training, Max Efficiency run with V100@160W | V100 performance measured on pre-production hardware.

SP ENERGY EFFICIENCY @ 28 NM

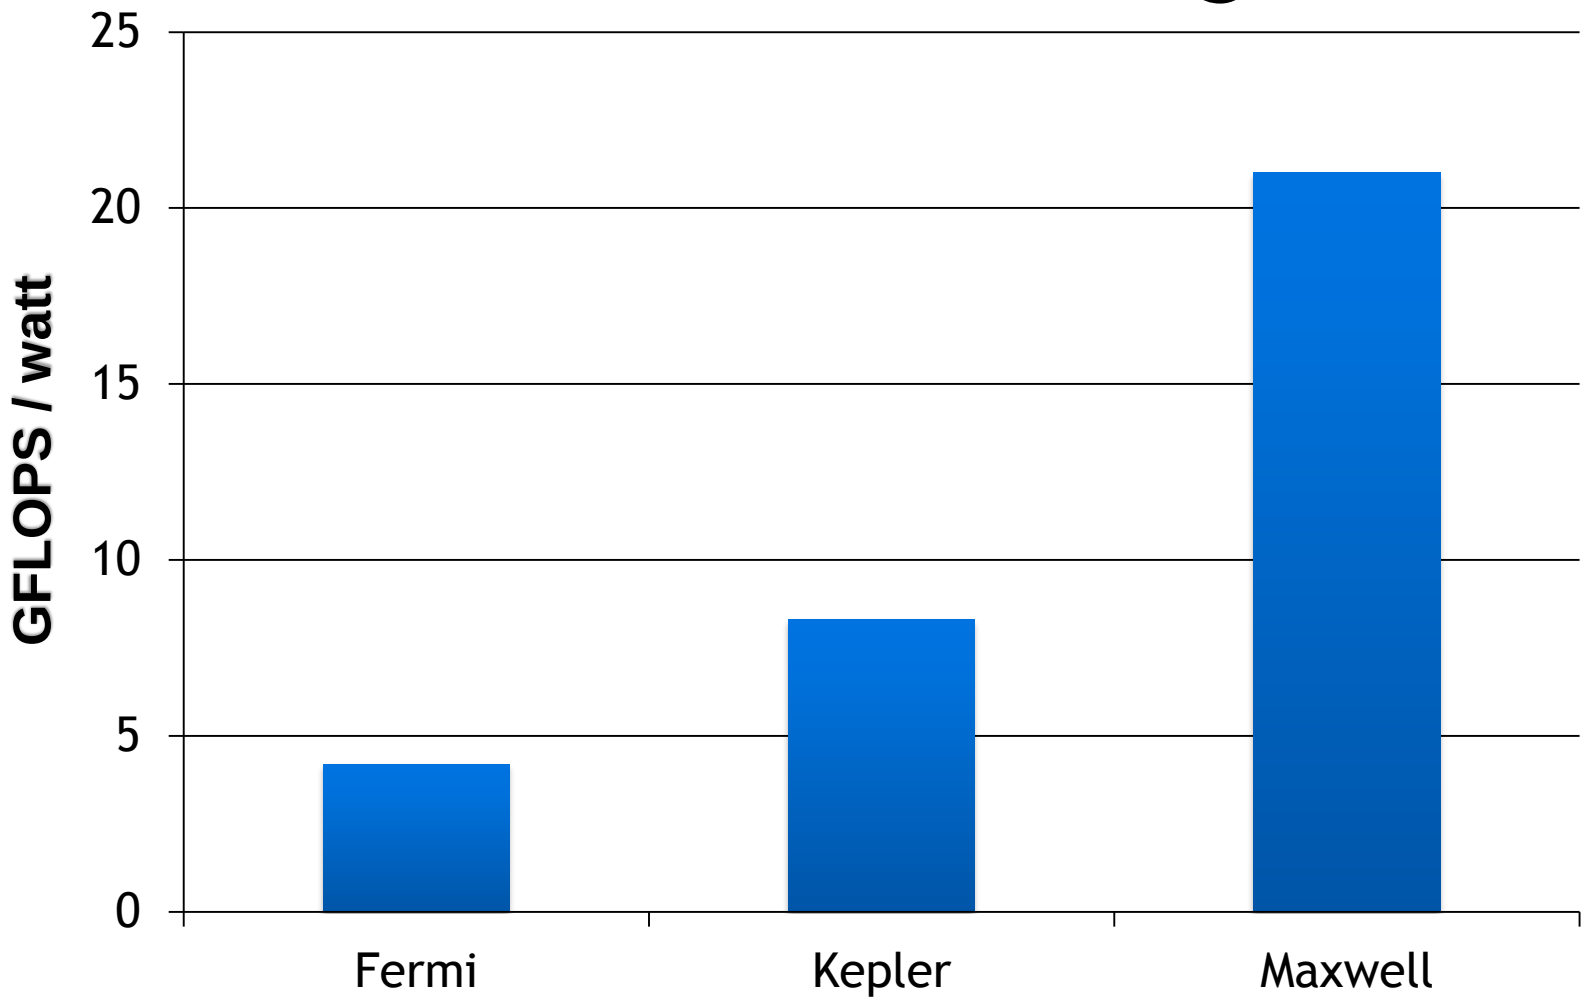

HETEROGENEOUS COMPUTING

OPTIMIZING SERIAL/PARALLEL EXECUTION

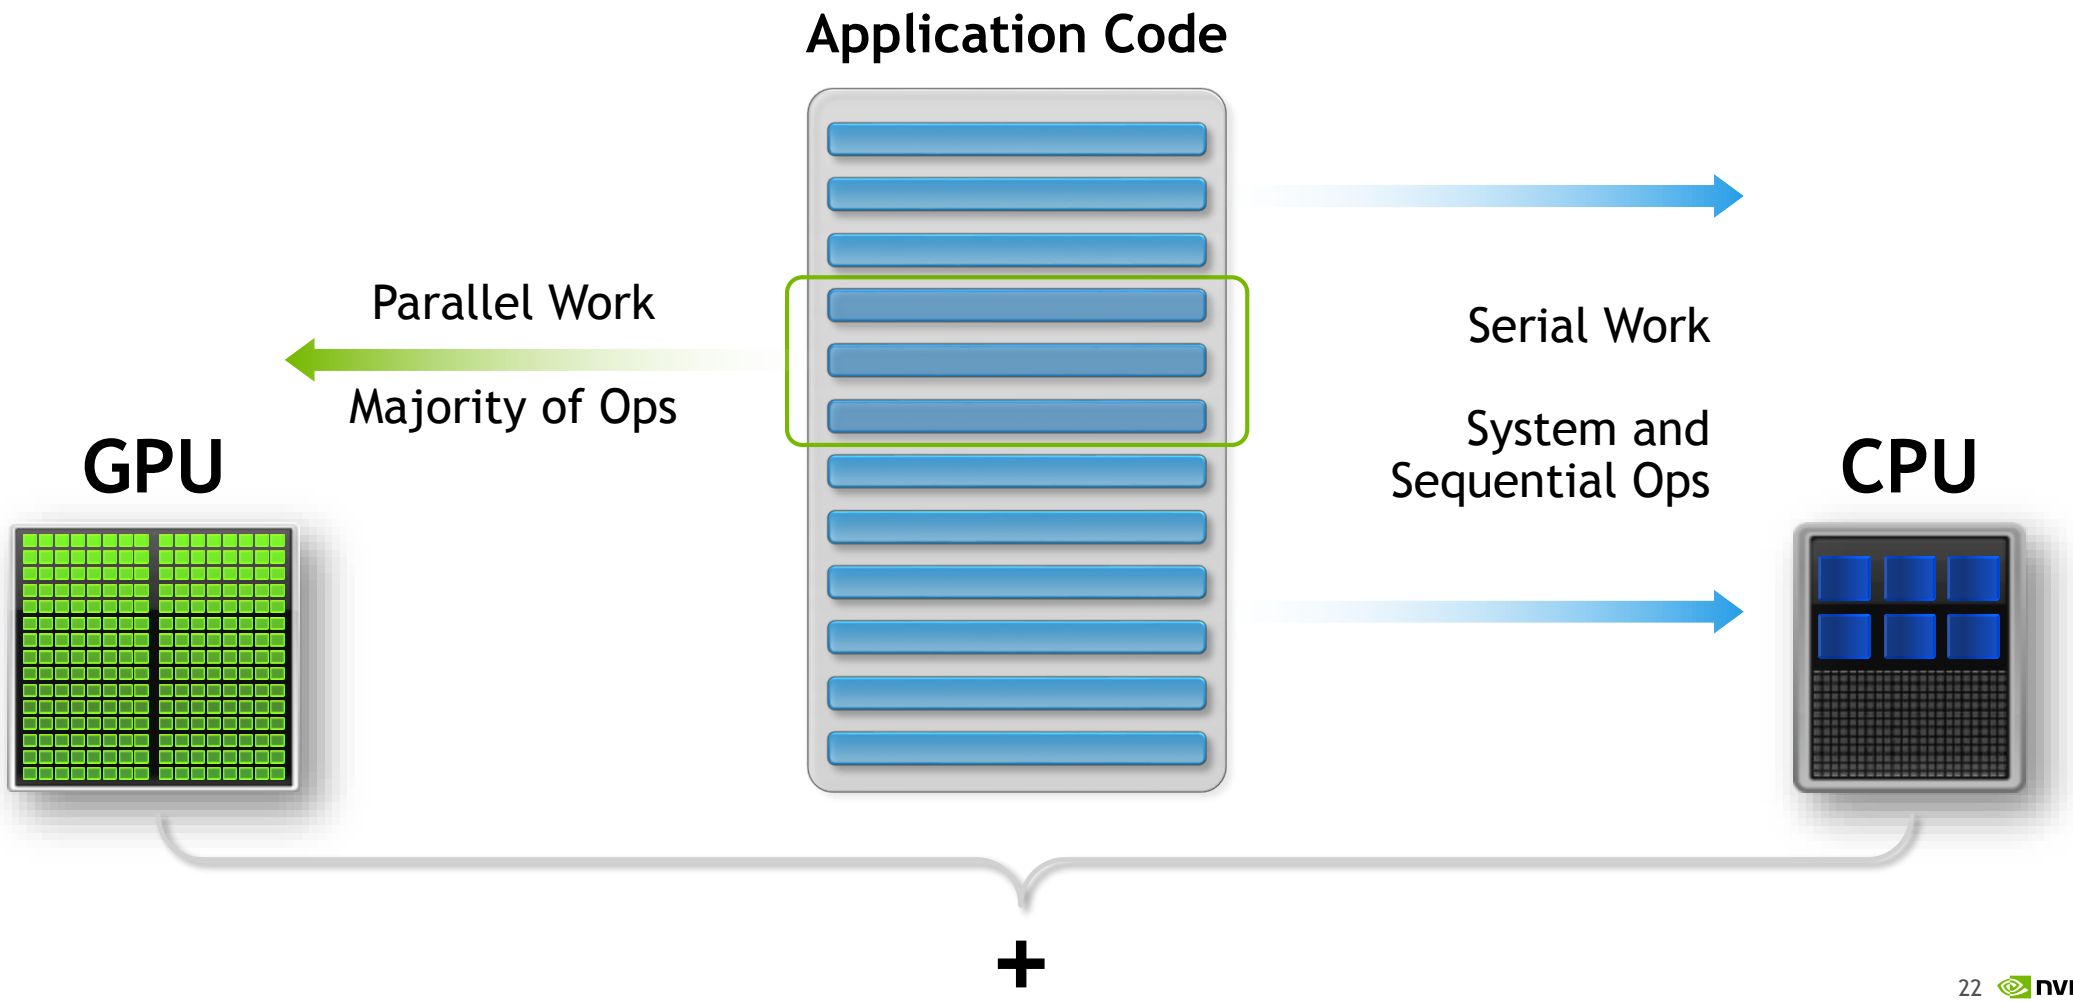

TWO TYPES OF ACCELERATORS

Many-Weak-Cores (MWC) Model

Single CPU Core for Both Serial & Parallel Work

Xeon Phi (*And Others*)
Many Weak Serial Cores

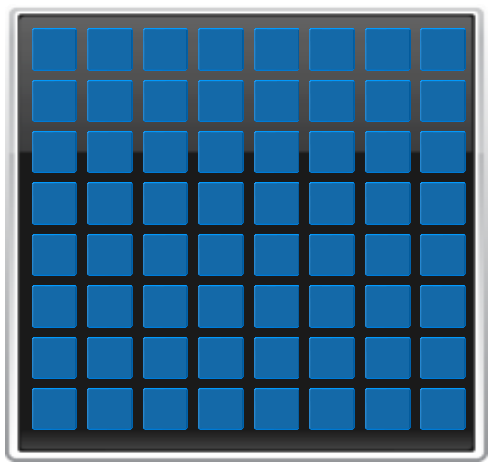

Heterogeneous Computing Model

Complementary Processors Work Together

CPU
Optimized for
Serial Tasks

GPU Accelerator
Optimized for
Parallel Tasks

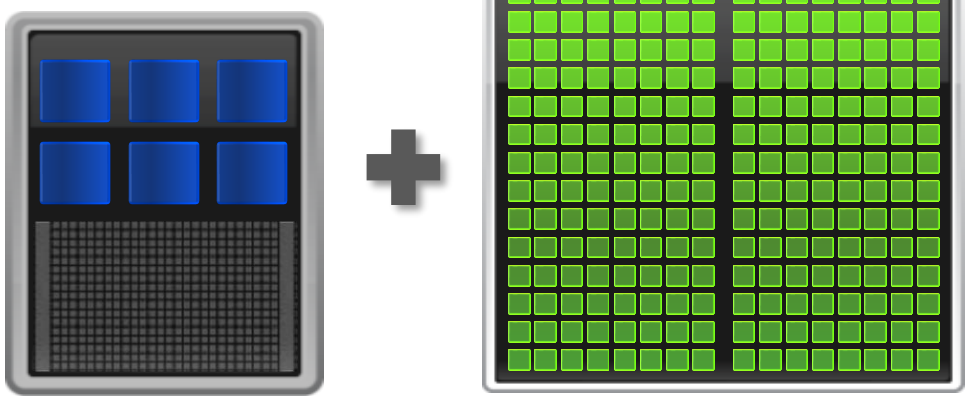

EXTENSIBILITY

NVLINK: A MEMORY FABRIC, NOT A NETWORK

DGX-1: 8 NVLink-Connected GPUs

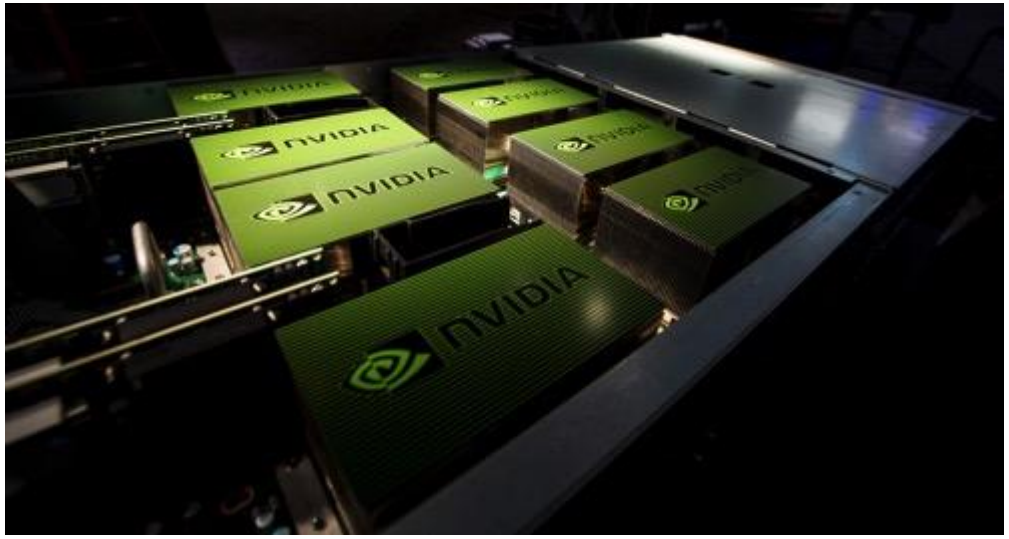

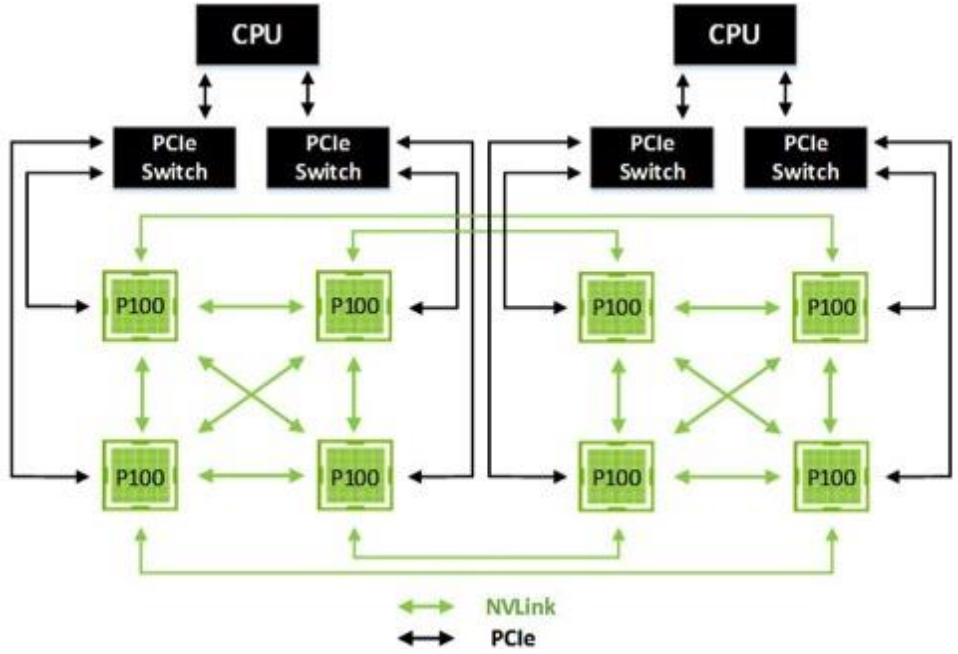

LATENCY HIDING FOR LOAD/STORE/ATOMICS

Where are the NICs? There are no NICs.

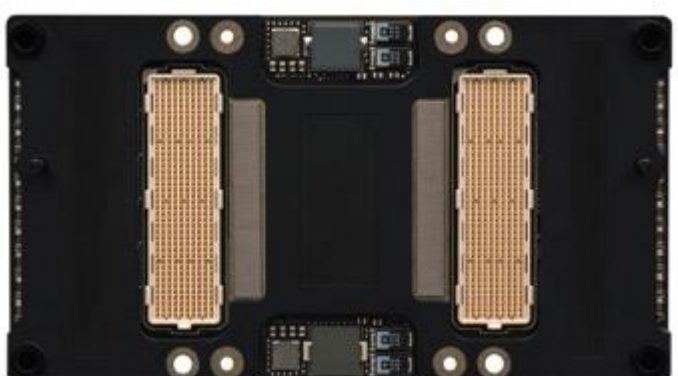

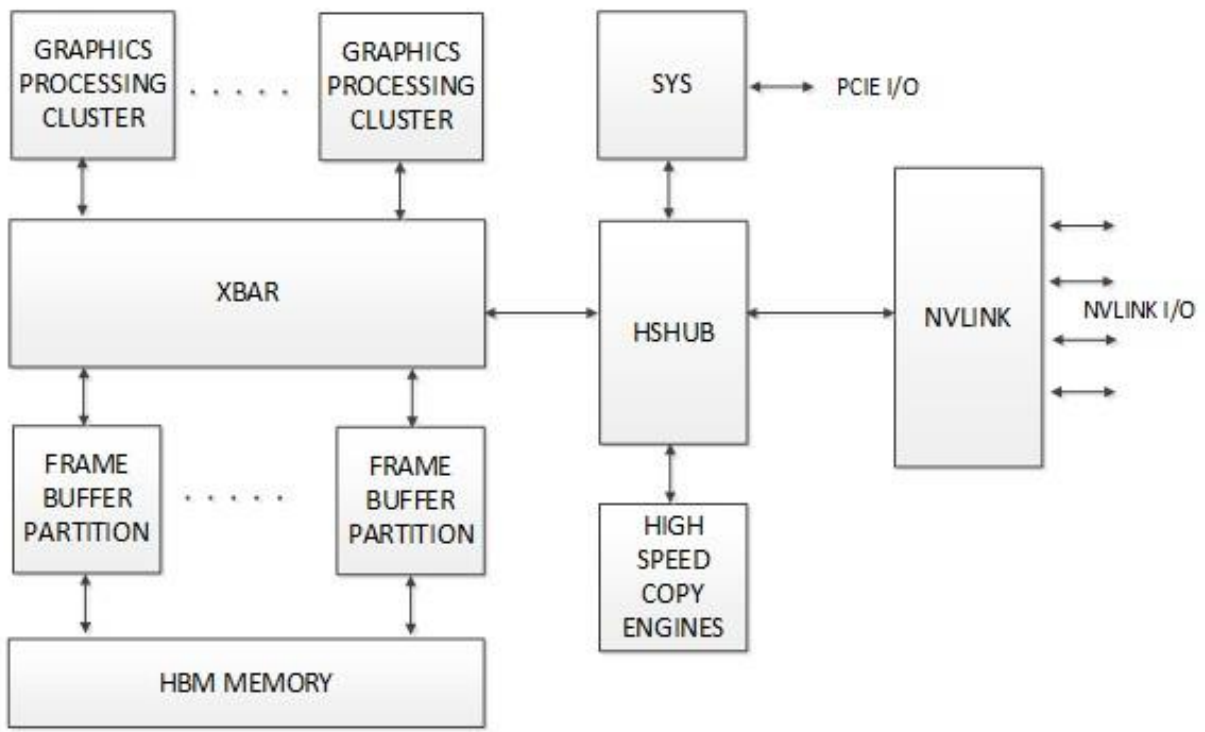

STRONG SCALING

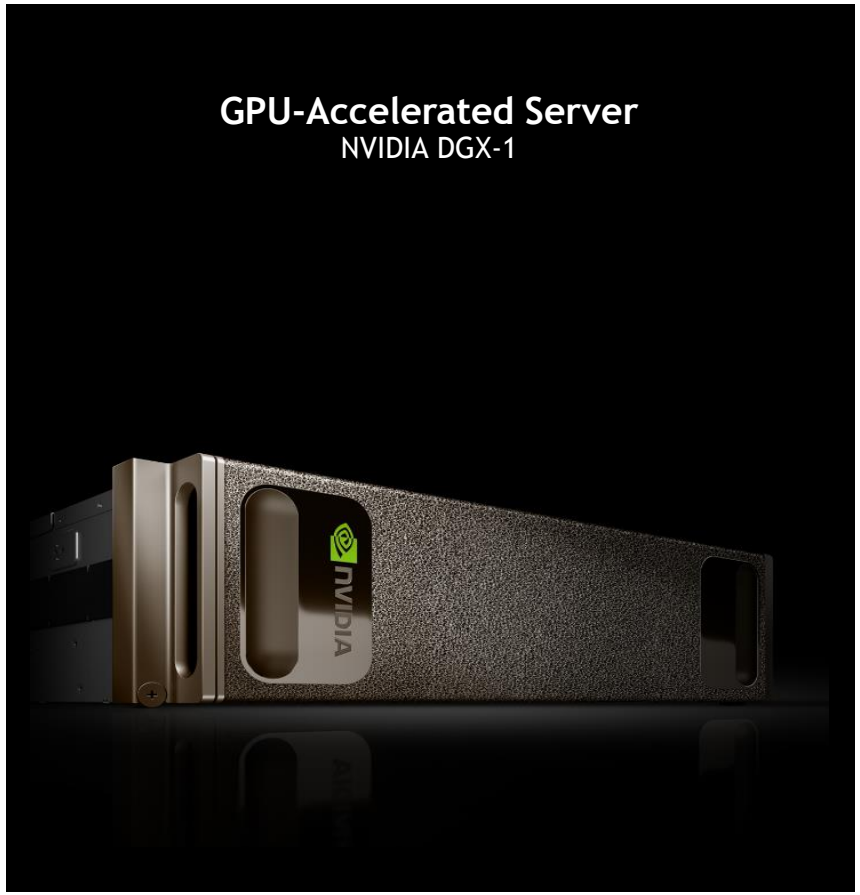

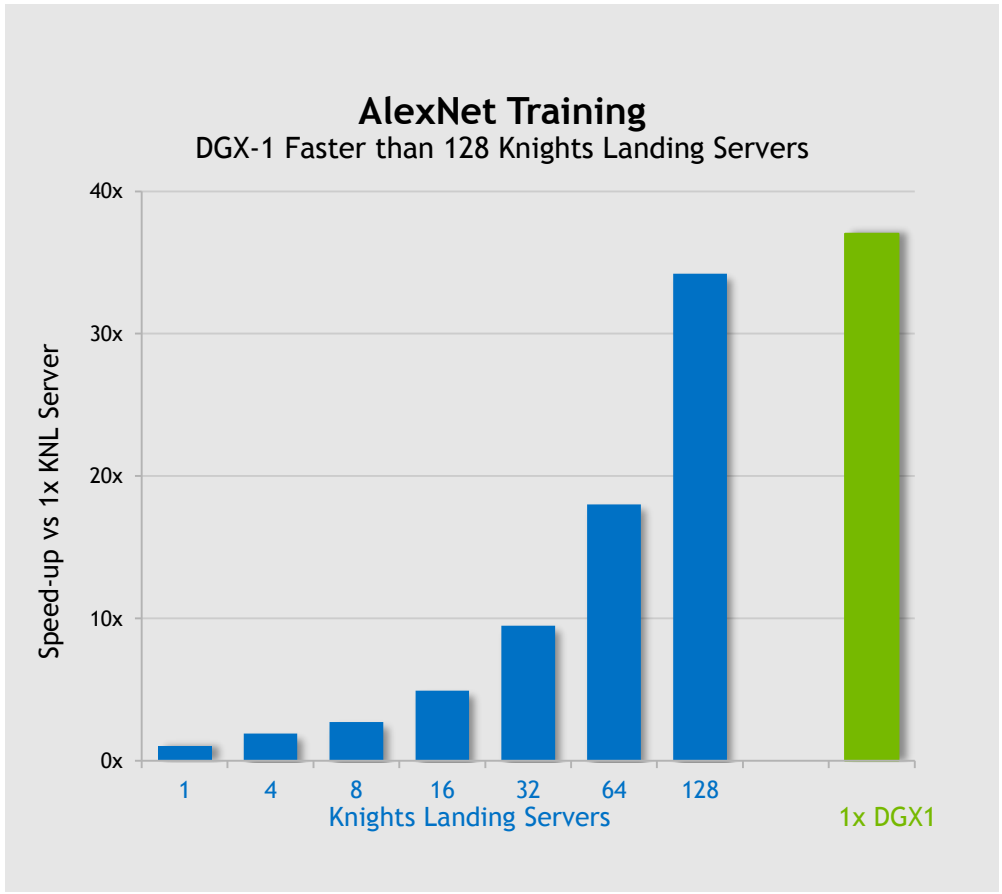

STRONG SCALING

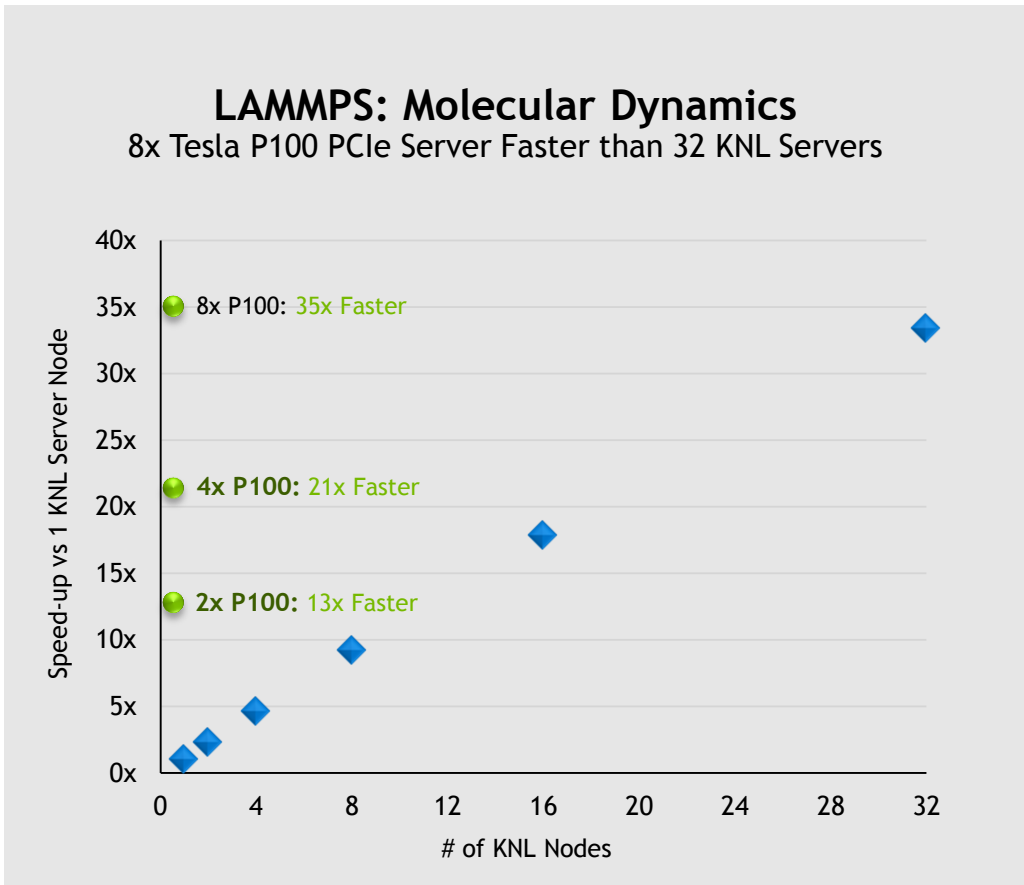

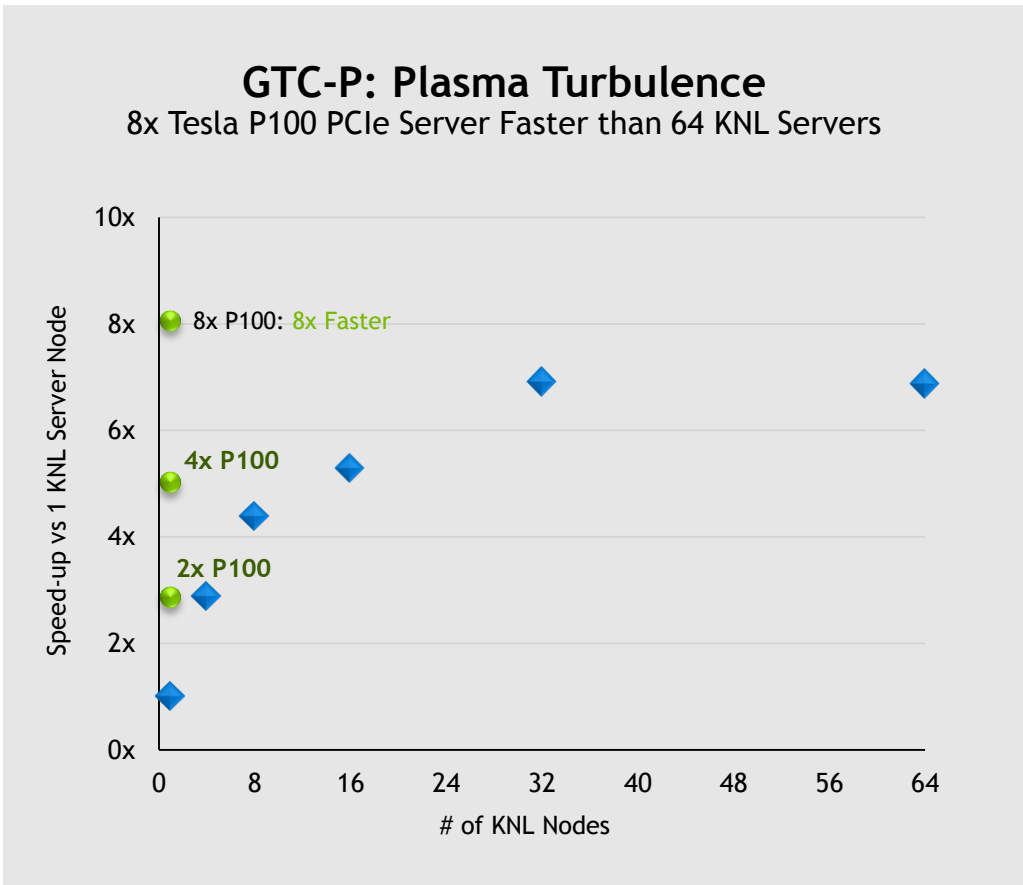

STRONG SCALING

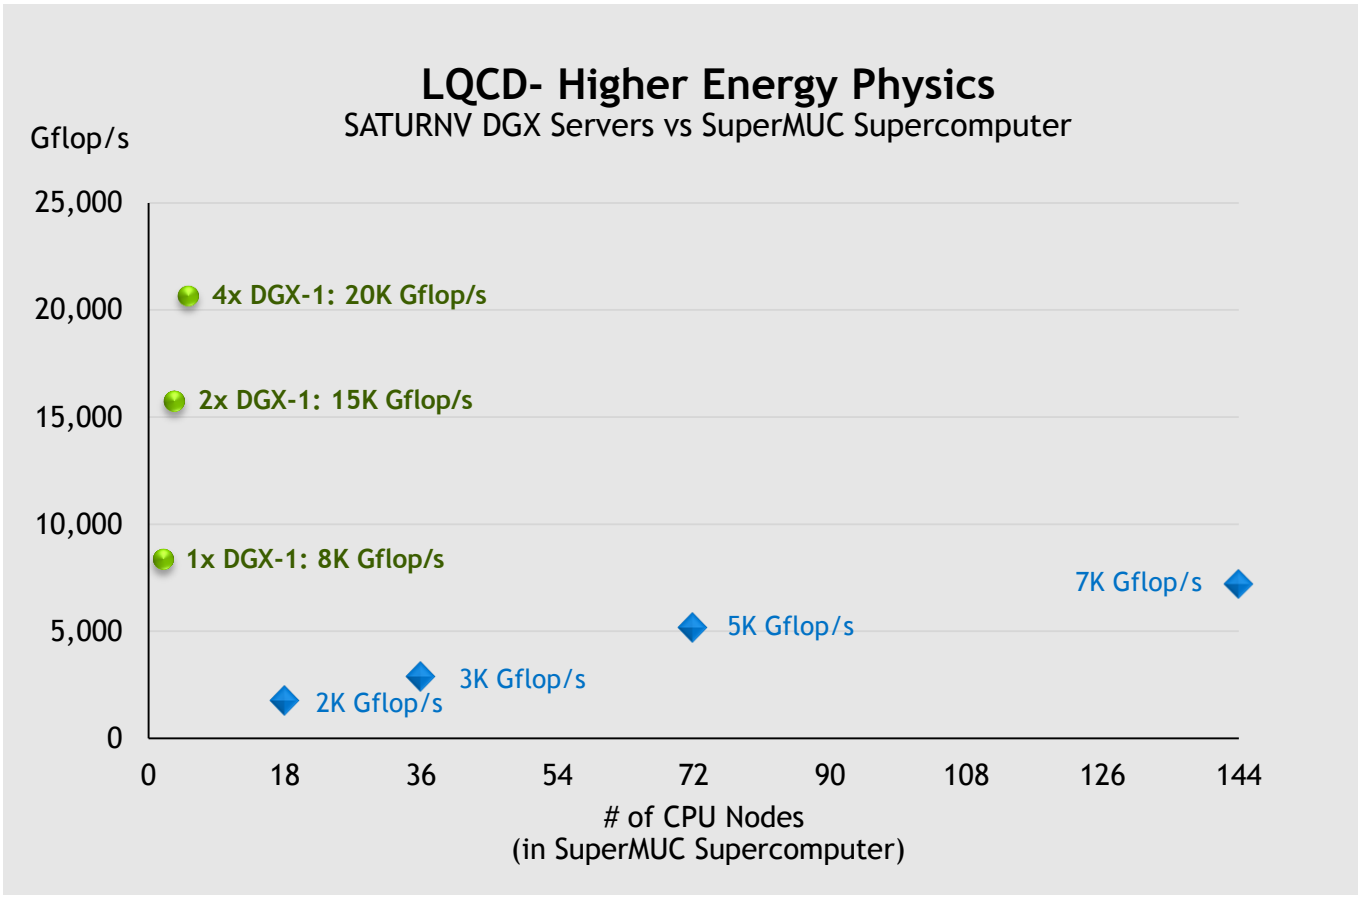

of CPU Servers to Match Performance of SATURNV

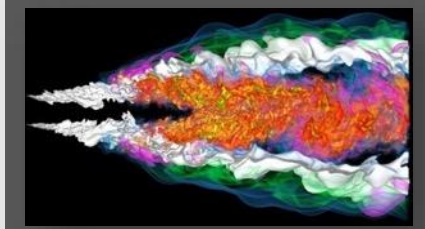

2,300
CPU Servers

S3D: Discovering New Fuel for Engines

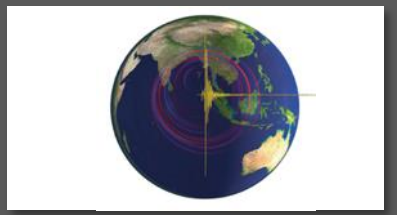

3,800
CPU Servers

SPECFEM3D: Simulating Earthquakes

QUDA version 0.9beta, using double-half mixed precision DDalphaAMG using double-single

NEW TENSOR CORE

New CUDA TensorOp instructions and data formats

4×4 matrix processing array

$D_{FP32} = A_{FP16} \times B_{FP16} + C_{FP32}$

Optimized for deep learning

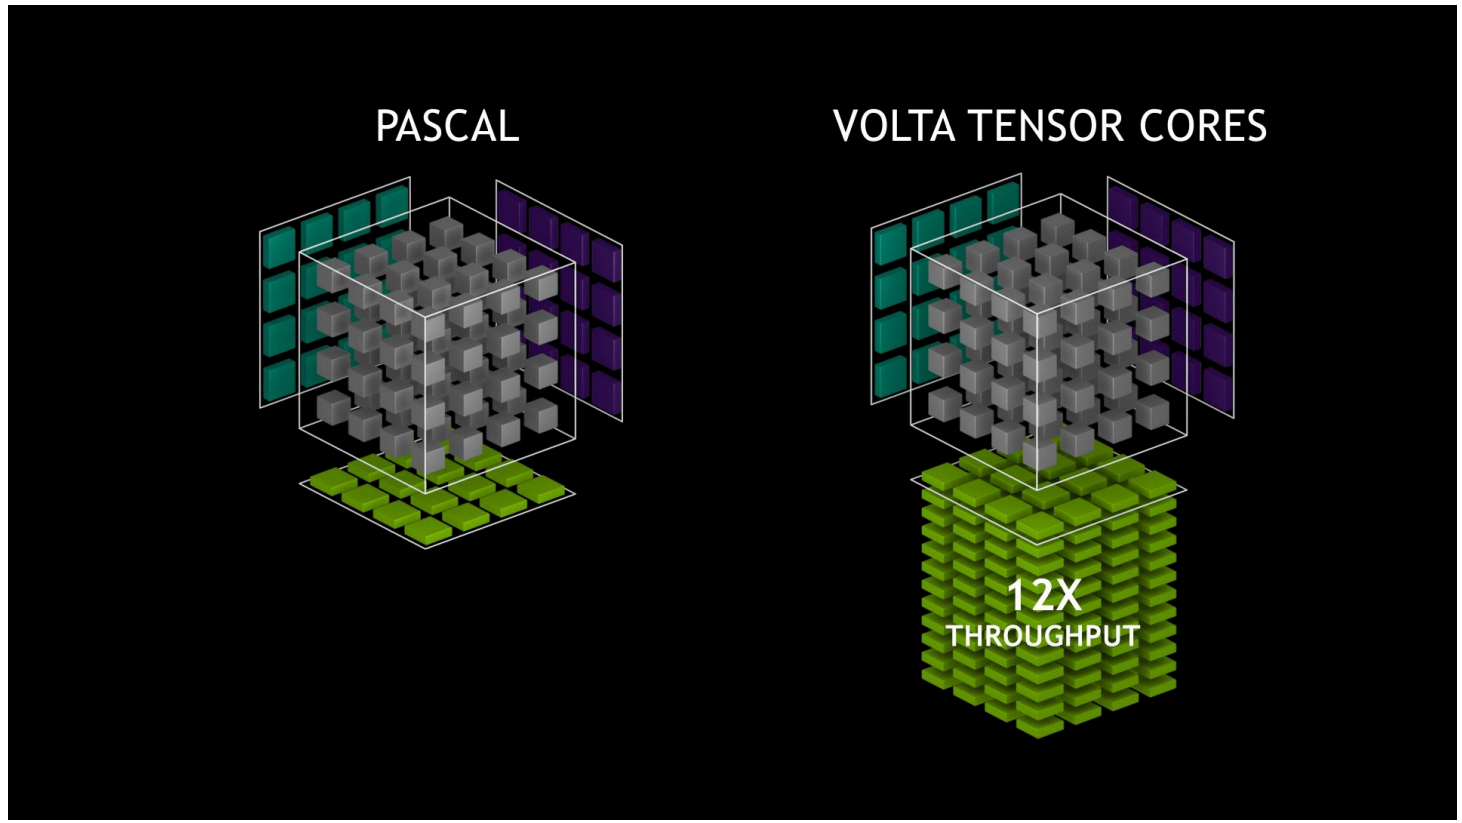

TESLA PLATFORM

TESLA IS A PLATFORM

World's Leading Data Center Platform for Accelerating HPC and AI

APPLICATIONS

INTERNET SERVICES

Alibaba, amazon, Baidu 百度, ebay, facebook, flickr, Google, Microsoft, Pinterest, skype, twitter, yelp

ENTERPRISE APPLICATIONS

Healthcare, Manufacturing, Finance, Automotive, Retail, Defense, ...

HPC

ANSYS, SIMULIA, GROMACS FAST. FLEXIBLE. FREE., +450 Applications

INDUSTRY FRAMEWORKS & TOOLS

FRAMEWORKS

Caffe2, Microsoft Cognitive Toolkit, mxnet, PYTORCH, TensorFlow, theano

ECOSYSTEM TOOLS

allinea, ArrayFire, CONTINUUM ANALYTICS, WOLFRAM, MathWorks

NVIDIA SDK

DEEP LEARNING SDK

cuDNN, TensorRT, NCCL, cuBLAS, cuSPARSE, DeepStream SDK

COMPUTEWORKS

CUDA C/C++, PGI OpenACC

TESLA GPU & SYSTEMS

TESLA GPU

NVIDIA DGX-1

NVIDIA HGX-1

DELL, Hewlett Packard Enterprise, IBM

SYSTEM OEM

amazon web services, Google Cloud Platform, Microsoft Azure

CLOUD

NVIDIA

MULTIPLE GROWTH MARKETS

GROWTH MARKETS

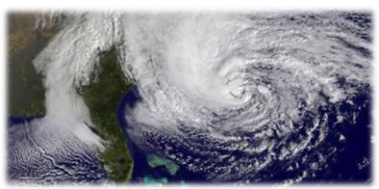

HPC

AI Training

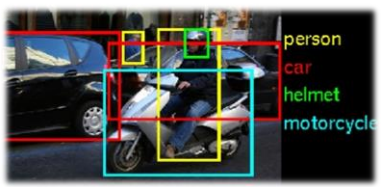

AI Inference

Desktop Virtualization

Video Transcoding

TESLA PLATFORM

CONCLUSION

PASCAL TO VOLTA

Architecture with Technology

- Area: ~600 mm^2 → ~800 mm^2 (~33% more area)
- Process: ~ small Pascal → Volta improvement (a few percent)
- Clocks: similar dynamic range, power limited
- Memory BW (sustained): 50% improvement
- Communications (NVLink): 160 GB/s → 300 GB/s (almost double!)
- AI (Tensor Cores): ~20 TFLOPS → 120 TFLOPS (~6x!)

REVOLUTIONARY PERFORMANCE FOR HPC AND AI

Single Platform For Data Science and Computation Science

1.5X HPC Performance In 1 Year

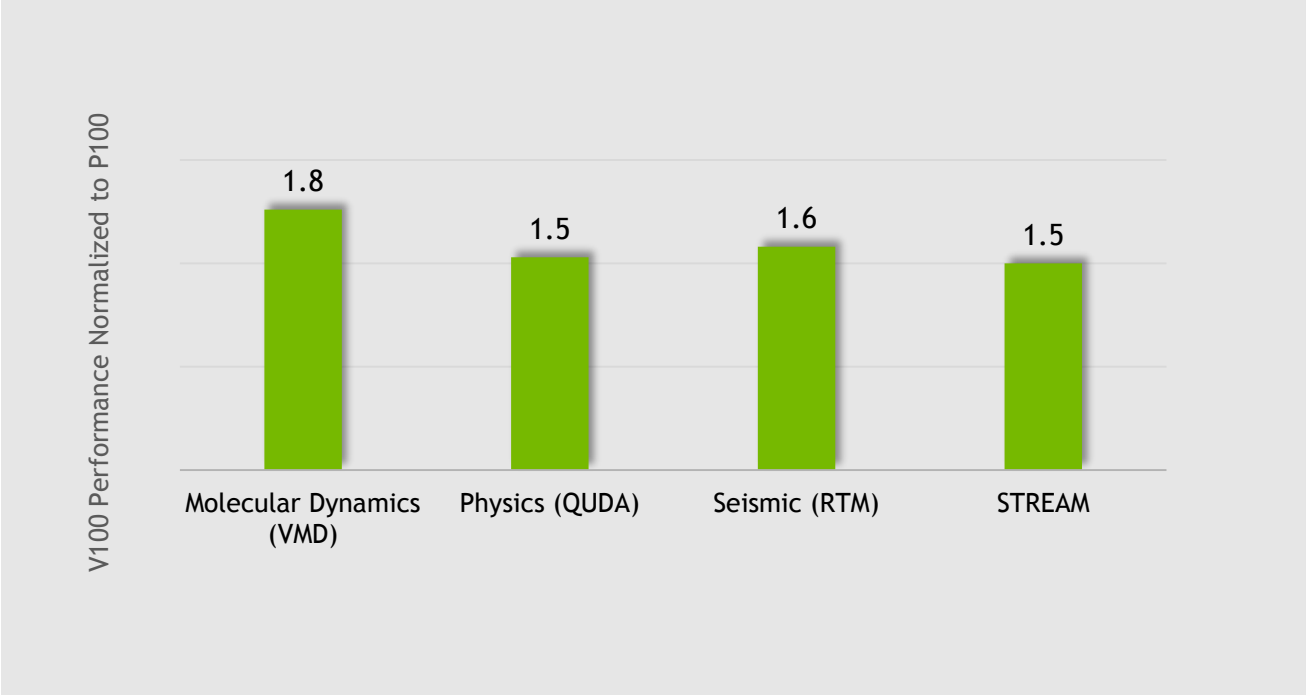

3X AI Performance in 1 Year

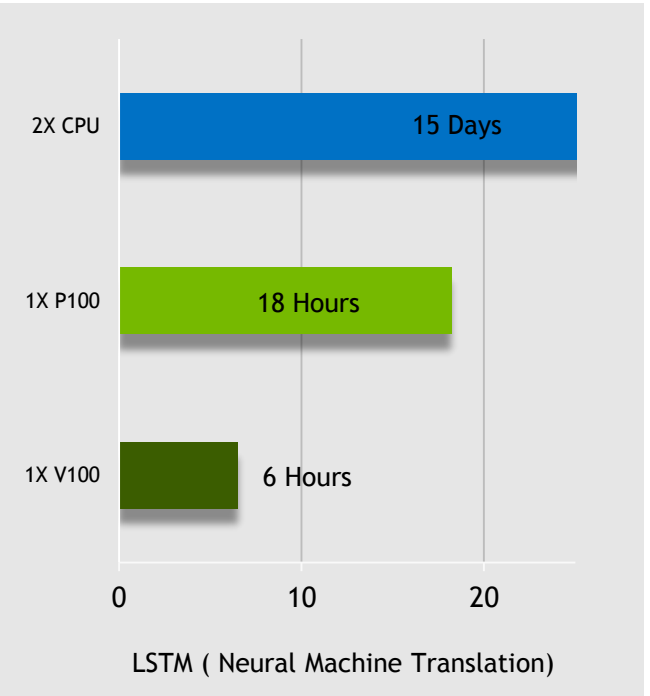

NMT Training for 13 Epochs |German ->English, WMT15 subset |System Config Info: CPU = 2x Xeon E5 2699 V4 w/ P100s or V100s |
QUDA, RTM, STREAM | System Config Info: 2X Xeon E5-2690 v4, 2.6GHz, w/ 1X Tesla P100 or V100.
VMD | System Config Info: Xeon E5-2698v3 w/ 1x Tesla P100 and E5-2697Av4 w/ 1x Tesla V100 | V100 measured on pre-production hardware.

GPU PERFORMANCE COMPARISON

	P100	V100	Ratio
Training acceleration	10 TOPS	120 TOPS	12x
Inference acceleration	21 TFLOPS	120 TOPS	6x
FP64/FP32	5/10 TFLOPS	7.5/15 TFLOPS	1.5x
HBM2 Bandwidth	720 GB/s	900 GB/s	1.2x
NVLink Bandwidth	160 GB/s	300 GB/s	1.9x
L2 Cache	4 MB	6 MB	1.5x
L1 Caches	1.3 MB	10 MB	7.7x

GPU TRAJECTORY

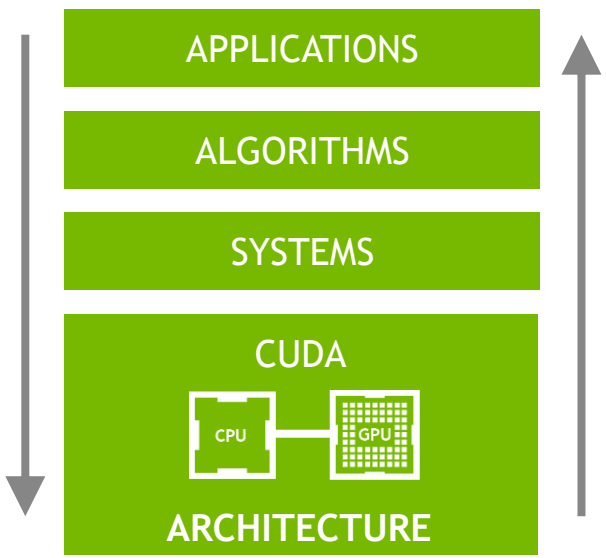

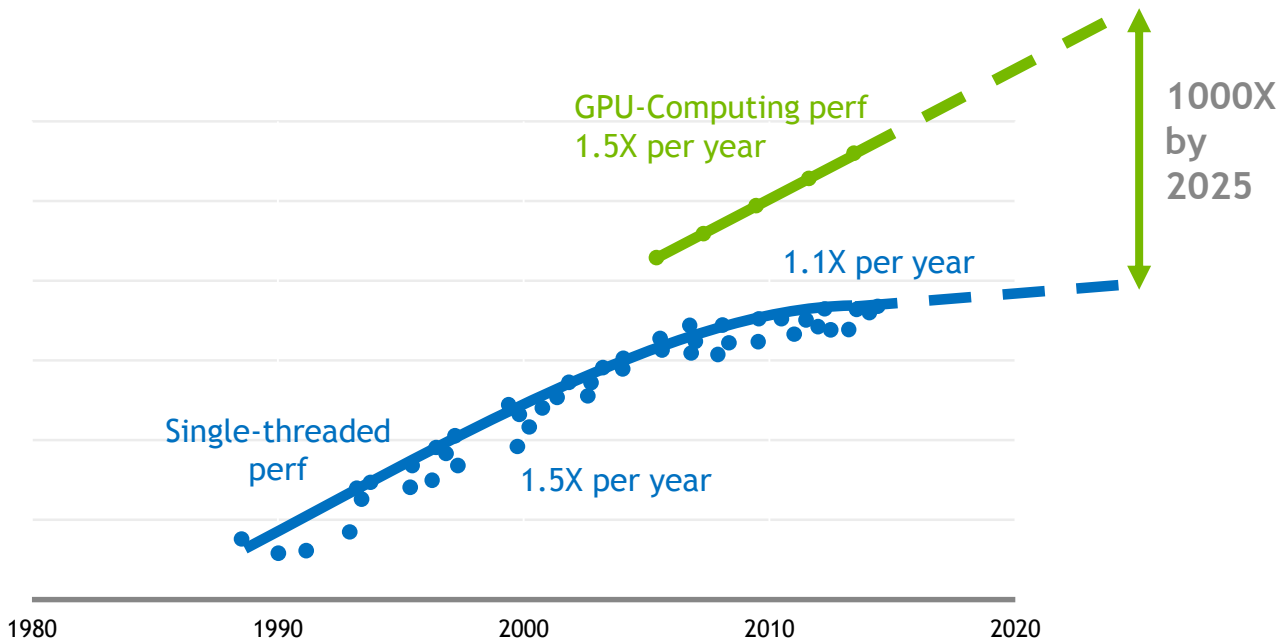

Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten New plot and data collected for 2010-2015 by K. Rupp

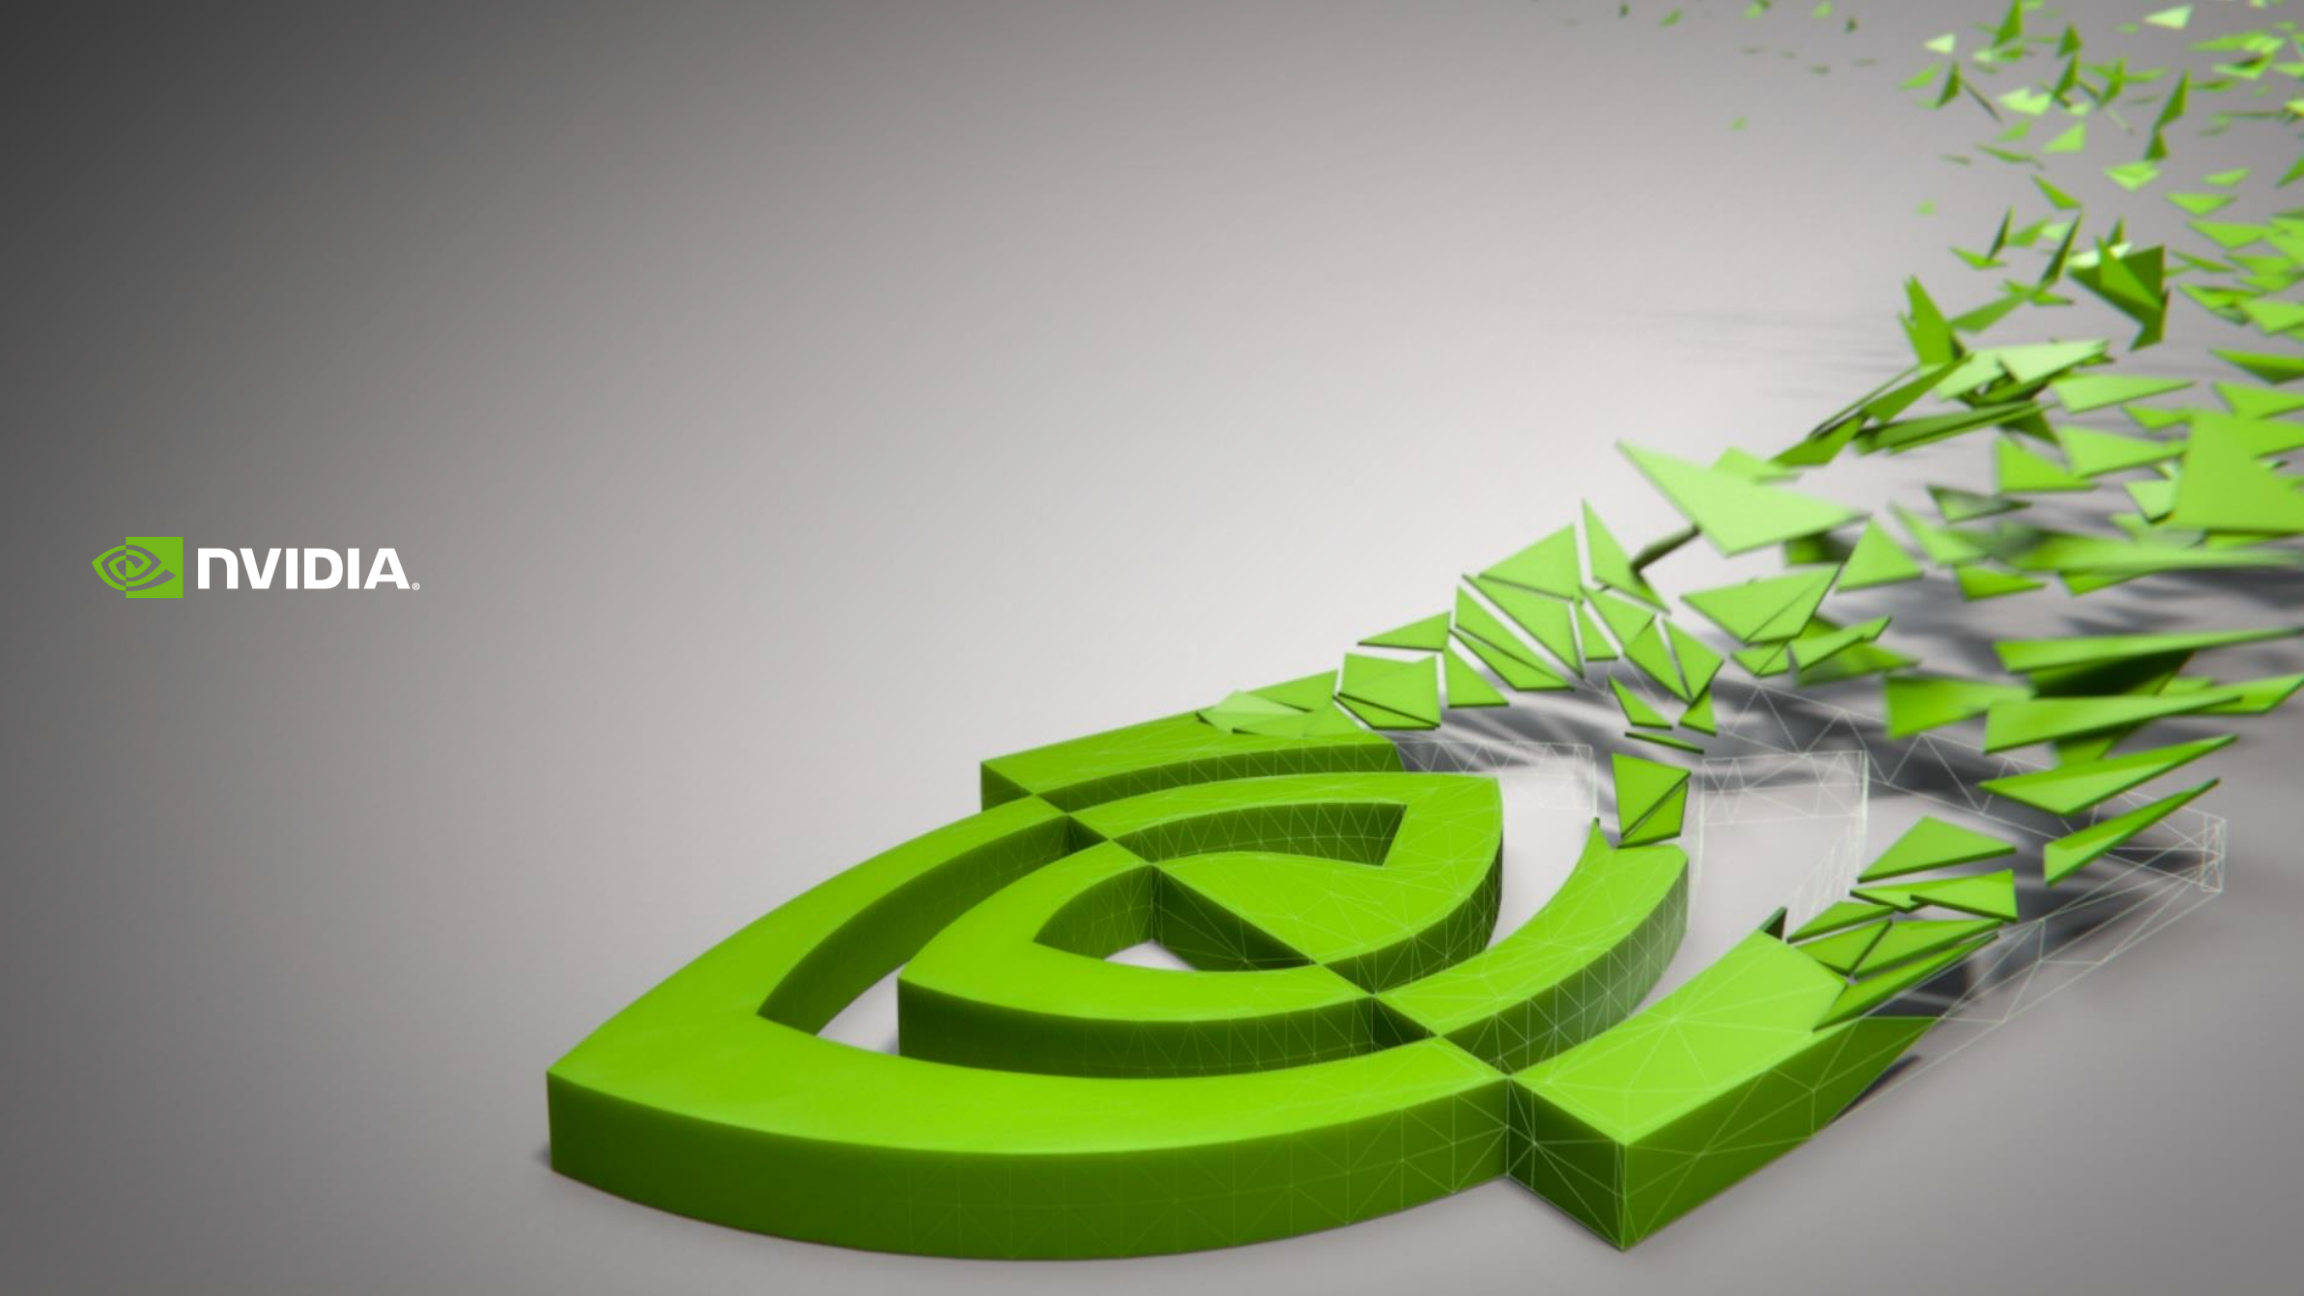
NVIDIA